



Introduction to Parallel & Distributed Computing

Lecture 4 Decomposition Techniques and OpenMP Task

Spring 2024

Instructor: 罗国杰

gluo@pku.edu.cn

Outline

◆ **Decomposition techniques**

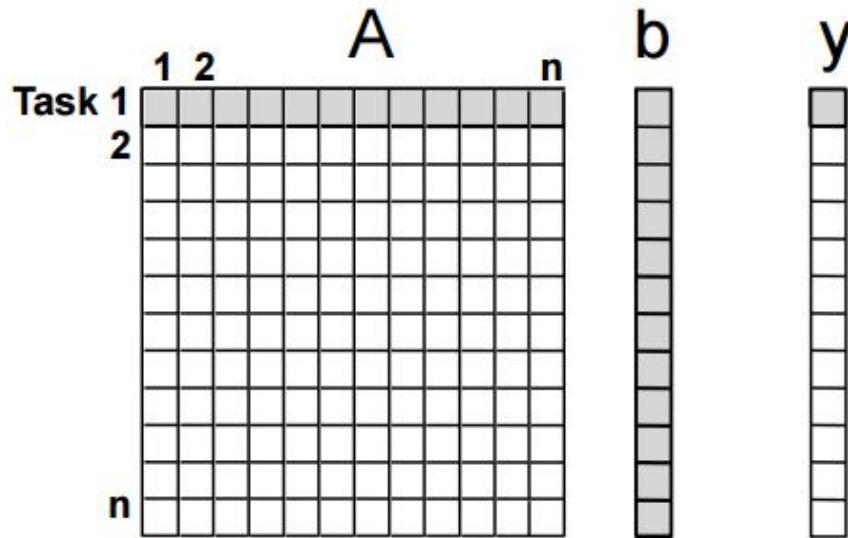
- **data decomposition**
- **recursive decomposition?**

◆ **OpenMP task**

◆ **Decomposition techniques (cont.)**

- **recursive decomposition**
- **exploratory decomposition**
- **speculative decomposition**

Example: Dense Matrix-Vector Product

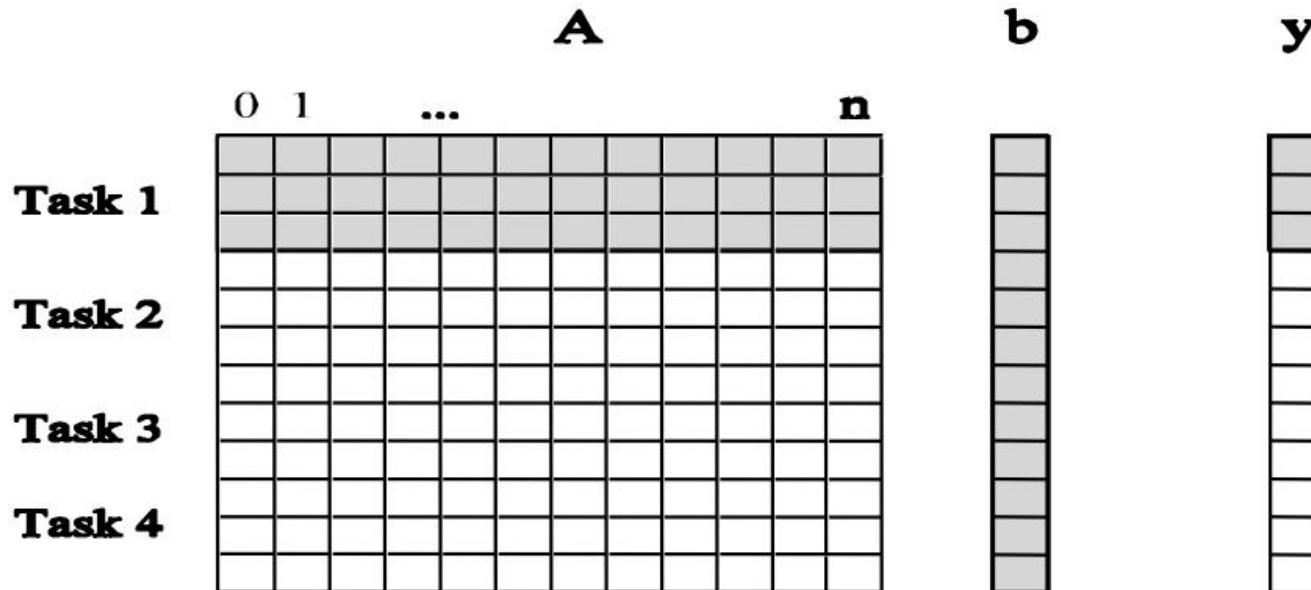


- ◆ Computing each element of output vector y is independent
- ◆ Easy to decompose dense matrix-vector product into tasks
 - one per element in y
- ◆ Observations
 - task size is uniform
 - no control dependences between tasks
 - tasks share b

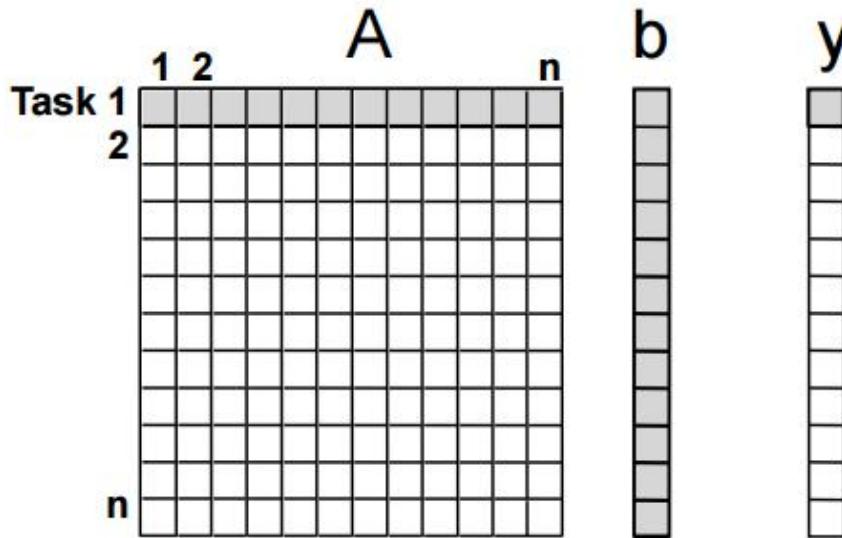
Source: COMP 422 at Rice Univ., Spring 2016.

Granularity of Decompositions

- ◆ Granularity = task size
 - depends on the number of tasks
- ◆ Fine-grain = large number of tasks
- ◆ Coarse-grain = small number of tasks
- ◆ Granularity examples for dense matrix-vector multiply
 - fine-grain: each task represents an individual element in y
 - coarser-grain: each task computes 3 elements in y



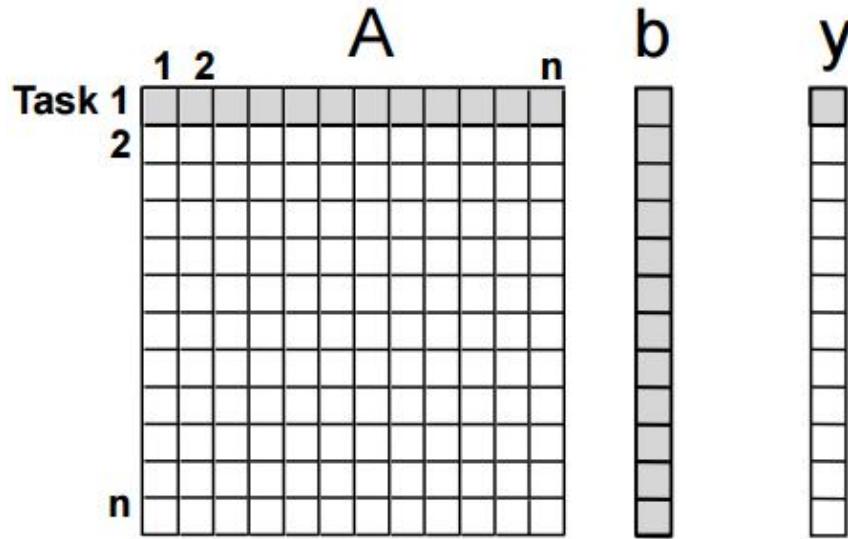
Example: Dense Matrix-Vector



- ◆ Computing each element of output vector y is independent
- ◆ Easy to decompose dense matrix-vector product into tasks
 - one per element in y
- ◆ Observations
 - task size is uniform
 - no control dependences between tasks
 - tasks share b

Question: Is n the maximum number of tasks possible?

Critical Path Length



Questions:

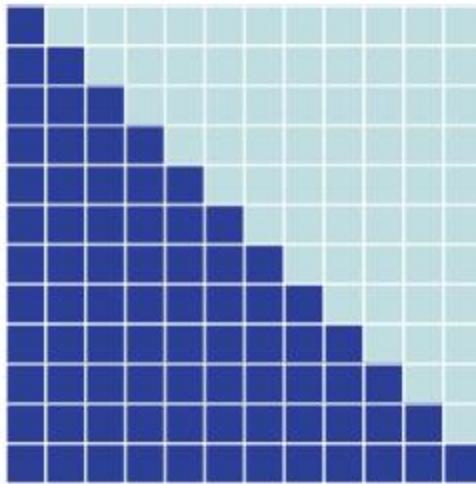
- What does a task dependency graph look like for DMVP?
- What is the shortest parallel execution time for the graph?
- How many processors are needed to achieve the minimum time?
- What is the maximum degree of concurrency?
- What is the average parallelism?

Exercise: Dense Matrix-Vector

- ◆ Try OpenMP with different granularities
- ◆ Which loop to parallelize?
- ◆ Related omp pragmas
 - parallel, for
- ◆ Related omp clauses
 - schedule

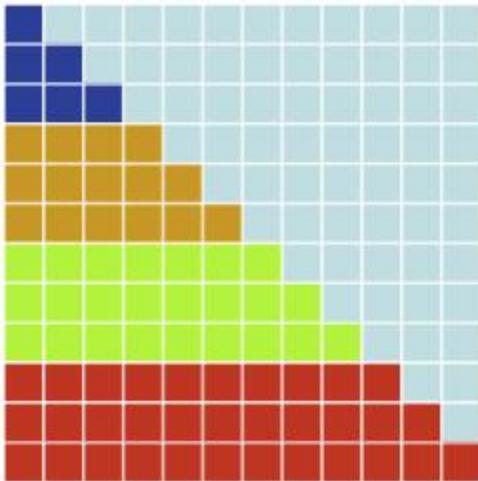
Loop Scheduling Example

```
for (int i = 0; i < 12; i++)  
    for (int j = 0; j <= i; j++)  
        a[i][j] = ...;
```



Loop Scheduling Example

```
#pragma omp parallel for  
for (int i = 0; i < 12; i++)  
    for (int j = 0; j <= i; j++)  
        a[i][j] = ...;
```



Typically, the iterations are divided by the number of threads and assigned as chunks to a thread

Loop Scheduling

◆ **Loop schedule**

- **How loop iterations are assigned to threads**
- **Static schedule**
 - Iterations assigned to threads before execution of loop
- **Dynamic schedule**
 - Iterations assigned to threads during execution of loop

◆ **The OpenMP *schedule* clause affects how loop iterations are mapped onto threads**

The schedule Clause

- ◆ *schedule (static [, chunk])*

- **Blocks of iterations of size “chunk” to threads**
 - **Round robin distribution**
 - **Low overhead, may cause load imbalance**

- ◆ **Best used for predictable and similar work per iteration**

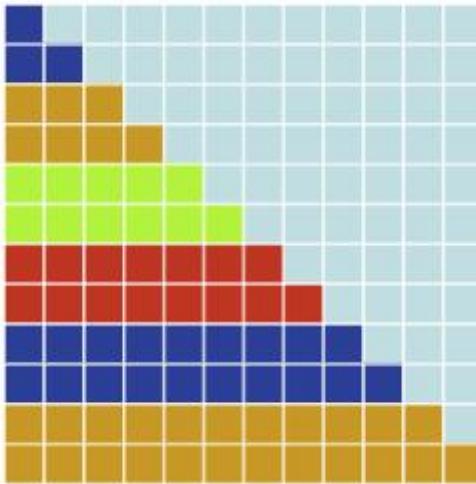
Loop Scheduling Example

```
#pragma omp parallel for schedule(static, 2)
```

```
for (int i = 0; i < 12; i++)
```

```
    for (int j = 0; j <= i; j++)
```

```
        a[i][j] = ...;
```



The schedule Clause

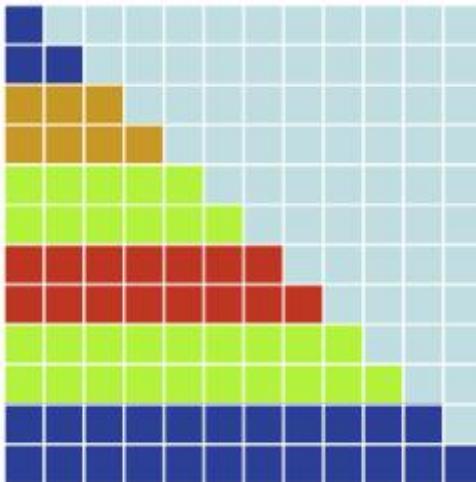
◆ *schedule (dynamic [, chunk])*

- **Threads grab “chunk” iterations**
- **When done with iterations, thread requests next set**
- **Higher threading overhead, can reduce load imbalance**

◆ **Best used for unpredictable or highly variable work**

Loop Scheduling Example

```
#pragma omp parallel for schedule(dynamic, 2)
for (int i = 0; i < 12; i++)
    for (int j = 0; j <= i; j++)
        a[i][j] = ...;
```



The schedule Clause

◆ *schedule (guided [, chunk])*

- **Dynamic schedule starting with large block**
- **Size of the blocks shrink; no smaller than “chunk”**
- **The initial block is proportional to**
 - `number_of_chunks / number_of_threads`
- **Subsequent blocks are proportional to**
 - `number_of_chunks_remaining / number_of_threads`
- ◆ **Best used as a special case of dynamic to reduce scheduling overhead when the computation gets progressively more time consuming**

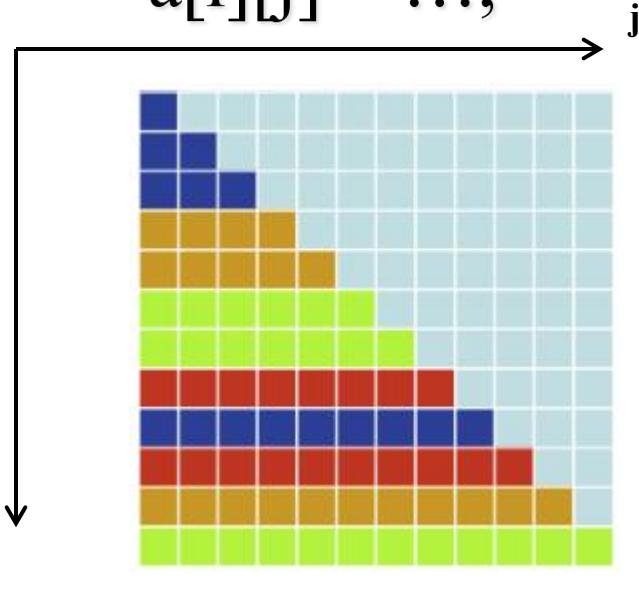
Loop Scheduling Example

```
#pragma omp parallel for schedule(guided)
```

```
for (int i = 0; i < 12; i++)
```

```
    for (int j = 0; j <= i; j++)
```

```
        a[i][j] = ...;
```



Decomposition Techniques

How should one decompose a task into various subtasks?

- ◆ **No single universal recipe**
- ◆ **In practice, a variety of techniques are used including**
 - **data decomposition**
 - **recursive decomposition**
 - **exploratory decomposition**
 - **speculative decomposition**

Data Decomposition

◆ **Steps**

1. identify the data on which computations are performed
2. partition the data across various tasks
 - partitioning induces a decomposition of the problem

◆ **Data can be partitioned in various ways**

- appropriate partitioning is critical to parallel performance

◆ **Decomposition based on**

- input data
- output data
- input + output data
- intermediate data

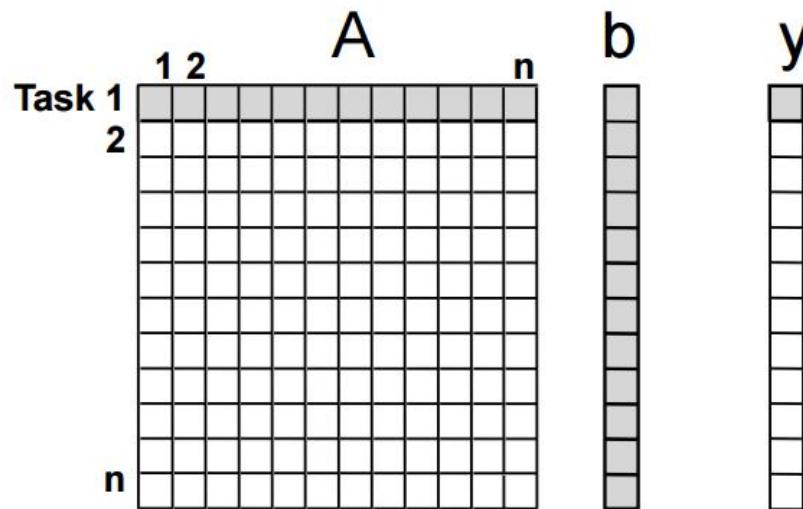
Decomposition Based on Input Data

- ◆ Applicable if each output is computed as a function of the input
- ◆ May be the only natural decomposition if output is unknown
 - examples
 - finding the minimum in a set or other reductions
 - sorting a vector
- ◆ Associate a task with each input data partition
 - task performs computation on its part of the data
 - subsequent processing combines partial results from earlier tasks

Decomposition Based on Output Data

- ◆ If each element of the output can be computed independently
- ◆ Partition the output data across tasks
- ◆ Have each task perform the computation for its outputs

Example:
dense matrix-vector
multiply



Output Data Decomposition: Example

- ◆ Matrix multiplication: $C = A \times B$
- ◆ Computation of C can be partitioned into four tasks

$$\begin{pmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{pmatrix} \cdot \begin{pmatrix} B_{1,1} & B_{1,2} \\ B_{2,1} & B_{2,2} \end{pmatrix} \rightarrow \begin{pmatrix} C_{1,1} & C_{1,2} \\ C_{2,1} & C_{2,2} \end{pmatrix}$$

Task 1: $C_{1,1} = A_{1,1}B_{1,1} + A_{1,2}B_{2,1}$

Task 2: $C_{1,2} = A_{1,1}B_{1,2} + A_{1,2}B_{2,2}$

Task 3: $C_{2,1} = A_{2,1}B_{1,1} + A_{2,2}B_{2,1}$

Task 4: $C_{2,2} = A_{2,1}B_{1,2} + A_{2,2}B_{2,2}$

Example: Decomposition Based on Output Data

- Count the frequency of item sets in database transactions

Database Transactions	Items	Itemset Frequency
A, B, C, E, G, H	A, B, C	1
B, D, E, F, K, L	D, E	3
A, B, F, H, L	C, F, G	0
D, E, F, H	A, E	2
F, G, H, K,	C, D	1
A, E, F, K, L	D, K	2
B, C, D, G, H, L	B, C, F	0
G, H, L	C, D, K	0
D, E, F, K, L		
F, G, H, L		

- Partition computation by partitioning the item sets to count
 - each task computes total count for each of its item sets

Database Transactions	Items	Itemset Frequency
A, B, C, E, G, H	A, B, C	1
B, D, E, F, K, L	D, E	3
A, B, F, H, L	C, F, G	0
D, E, F, H	A, E	2
F, G, H, K,		
A, E, F, K, L		
B, C, D, G, H, L		
G, H, L		
D, E, F, K, L		
F, G, H, L		

task 1

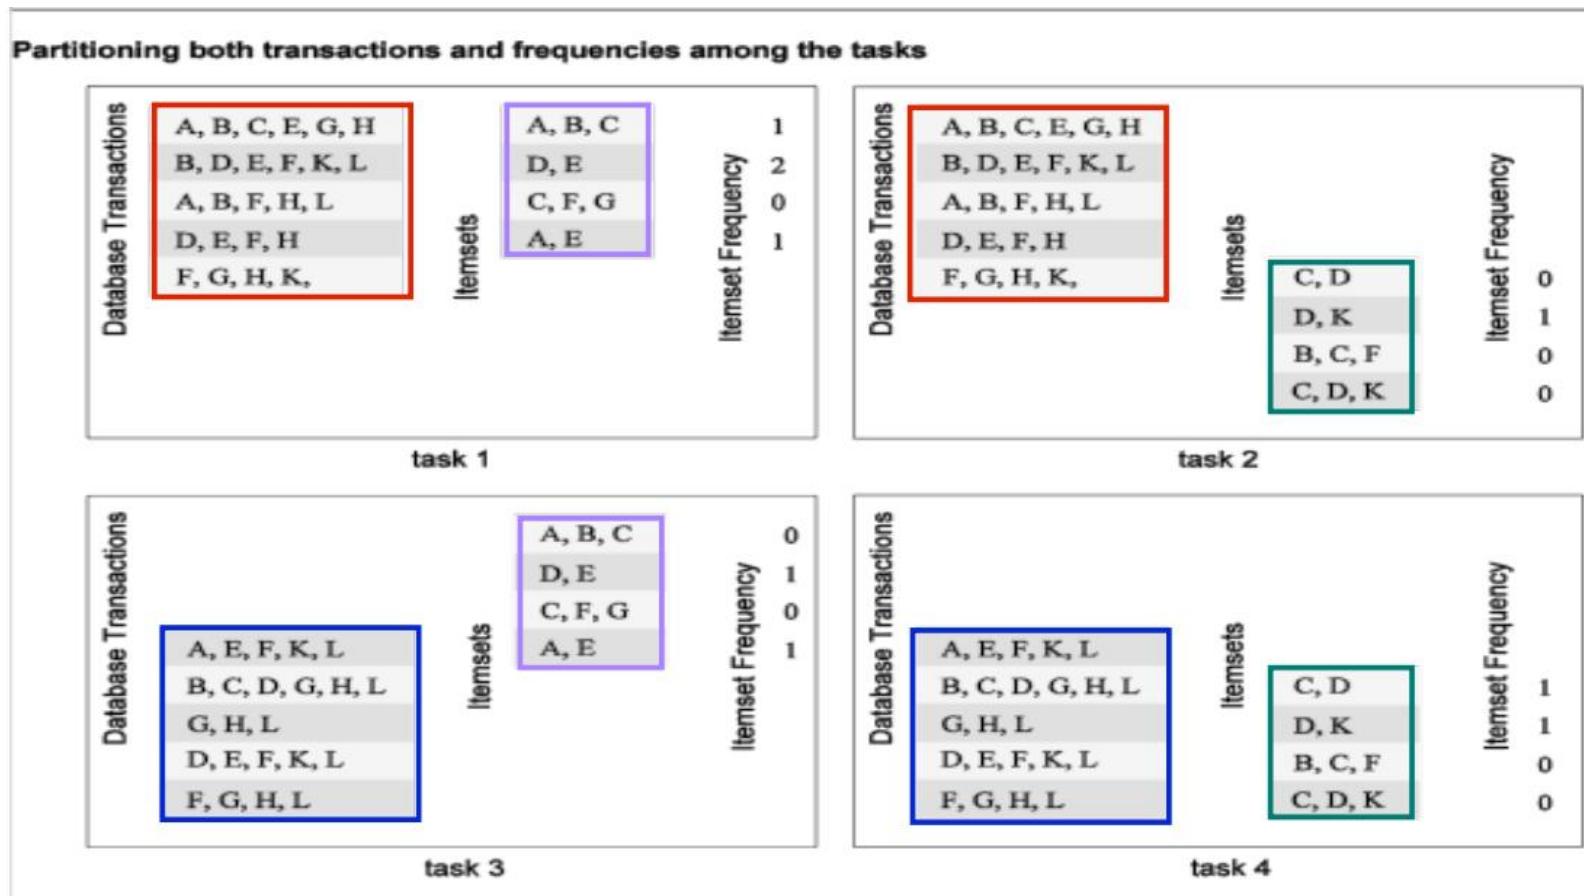
Database Transactions	Items	Itemset Frequency
A, B, C, E, G, H	A, B, C	1
B, D, E, F, K, L	D, K	2
A, B, F, H, L	B, C, F	0
D, E, F, H	C, D, K	0
F, G, H, K,		
A, E, F, K, L		
B, C, D, G, H, L		
G, H, L		
D, E, F, K, L		
F, G, H, L		

task 2

- append total counts for item subsets to produce result

Partitioning Input and Output Data

- ◆ Partition on both input and output for more concurrency
- ◆ Example: item set counting



Source: COMP 422 at Rice Univ., Spring 2016.

Intermediate Data Partitioning

- ◆ If computation is a sequence of transforms
 - (from input data to output data)
- ◆ Can decompose based on data for intermediate stages

Example: Intermediate Data Partitioning

- ◆ Dense Matrix Multiply
- ◆ Decomposition of intermediate data: yields $8 + 4$ tasks

Stage I

$$\begin{pmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{pmatrix} \cdot \begin{pmatrix} B_{1,1} & B_{1,2} \\ B_{2,1} & B_{2,2} \end{pmatrix} \rightarrow \begin{pmatrix} \left(\begin{array}{cc} D_{1,1,1} & D_{1,1,2} \\ D_{1,2,1} & D_{1,2,2} \\ D_{2,1,1} & D_{2,1,2} \\ D_{2,2,1} & D_{2,2,2} \end{array} \right) \end{pmatrix} \end{pmatrix}$$

Stage II

$$\begin{pmatrix} D_{1,1,1} & D_{1,1,2} \\ D_{1,2,1} & D_{1,2,2} \end{pmatrix} + \begin{pmatrix} D_{2,1,1} & D_{2,1,2} \\ D_{2,2,1} & D_{2,2,2} \end{pmatrix} \rightarrow \begin{pmatrix} C_{1,1} & C_{1,2} \\ C_{2,1} & C_{2,2} \end{pmatrix}$$

Task 01: $D_{1,1,1} = A_{1,1} B_{1,1}$

Task 02: $D_{2,1,1} = A_{1,2} B_{2,1}$

Task 03: $D_{1,1,2} = A_{1,1} B_{1,2}$

Task 04: $D_{2,1,2} = A_{1,2} B_{2,2}$

Task 05: $D_{1,2,1} = A_{2,1} B_{1,1}$

Task 06: $D_{2,2,1} = A_{2,2} B_{2,1}$

Task 07: $D_{1,2,2} = A_{2,1} B_{1,2}$

Task 08: $D_{2,2,2} = A_{2,2} B_{2,2}$

Task 09: $C_{1,1} = D_{1,1,1} + D_{2,1,1}$

Task 10: $C_{1,2} = D_{1,1,2} + D_{2,1,2}$

Task 11: $C_{2,1} = D_{1,2,1} + D_{2,2,1}$

Task 12: $C_{2,2} = D_{1,2,2} + D_{2,2,2}$

Intermediate Data Partitioning: Example

Tasks: dense matrix multiply decomposition of intermediate data

Task 01: $D_{1,1,1} = A_{1,1} B_{1,1}$

Task 03: $D_{1,1,2} = A_{1,1} B_{1,2}$

Task 05: $D_{1,2,1} = A_{2,1} B_{1,1}$

Task 07: $D_{1,2,2} = A_{2,1} B_{1,2}$

Task 09: $C_{1,1} = D_{1,1,1} + D_{2,1,1}$

Task 11: $C_{2,1} = D_{1,2,1} + D_{2,2,1}$

Task 02: $D_{2,1,1} = A_{1,2} B_{2,1}$

Task 04: $D_{2,1,2} = A_{1,2} B_{2,2}$

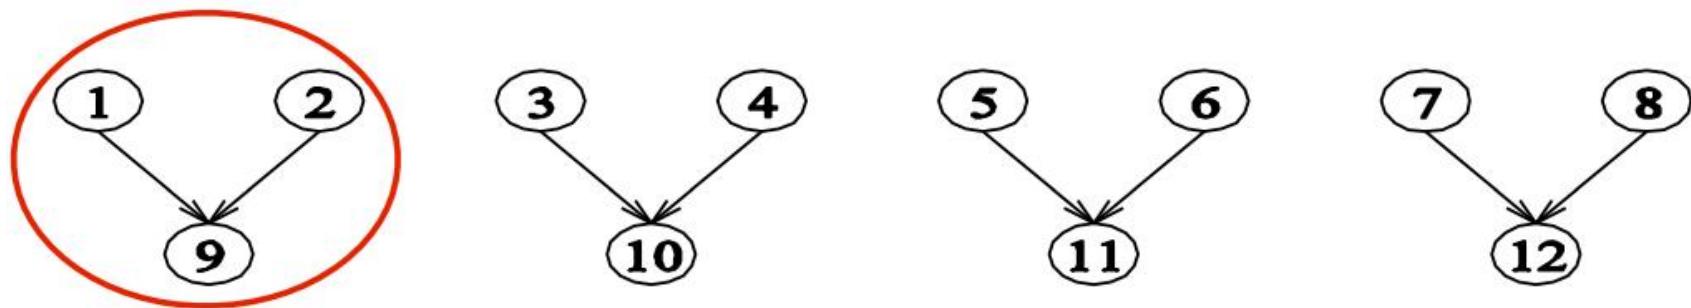
Task 06: $D_{2,2,1} = A_{2,2} B_{2,1}$

Task 08: $D_{2,2,2} = A_{2,2} B_{2,2}$

Task 10: $C_{1,2} = D_{1,1,2} + D_{2,1,2}$

Task 12: $C_{2,2} = D_{1,2,2} + D_{2,2,2}$

Task dependency graph



Owner Computes Rule

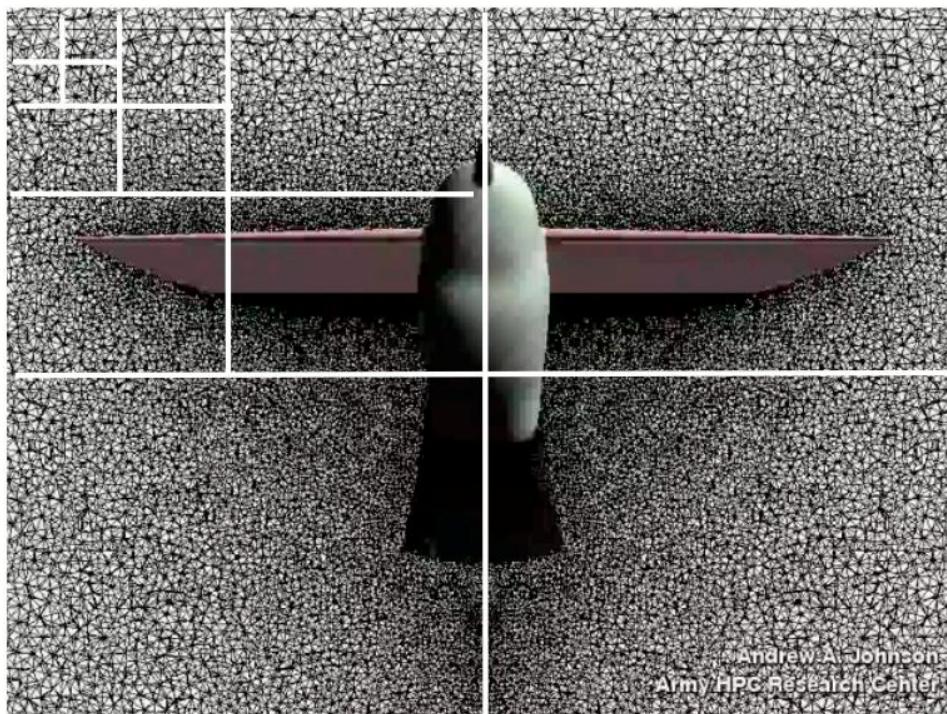
- ◆ **Each datum is assigned to a thread**
- ◆ **Each thread computes values associated with its data**
- ◆ **Implications**
 - **input data decomposition**
 - all computations using an input datum are performed by its thread
 - **output data decomposition**
 - an output is computed by the thread assigned to the output data

Recursive Decomposition

Suitable for problems solvable using divide-and-conquer 分治

Steps

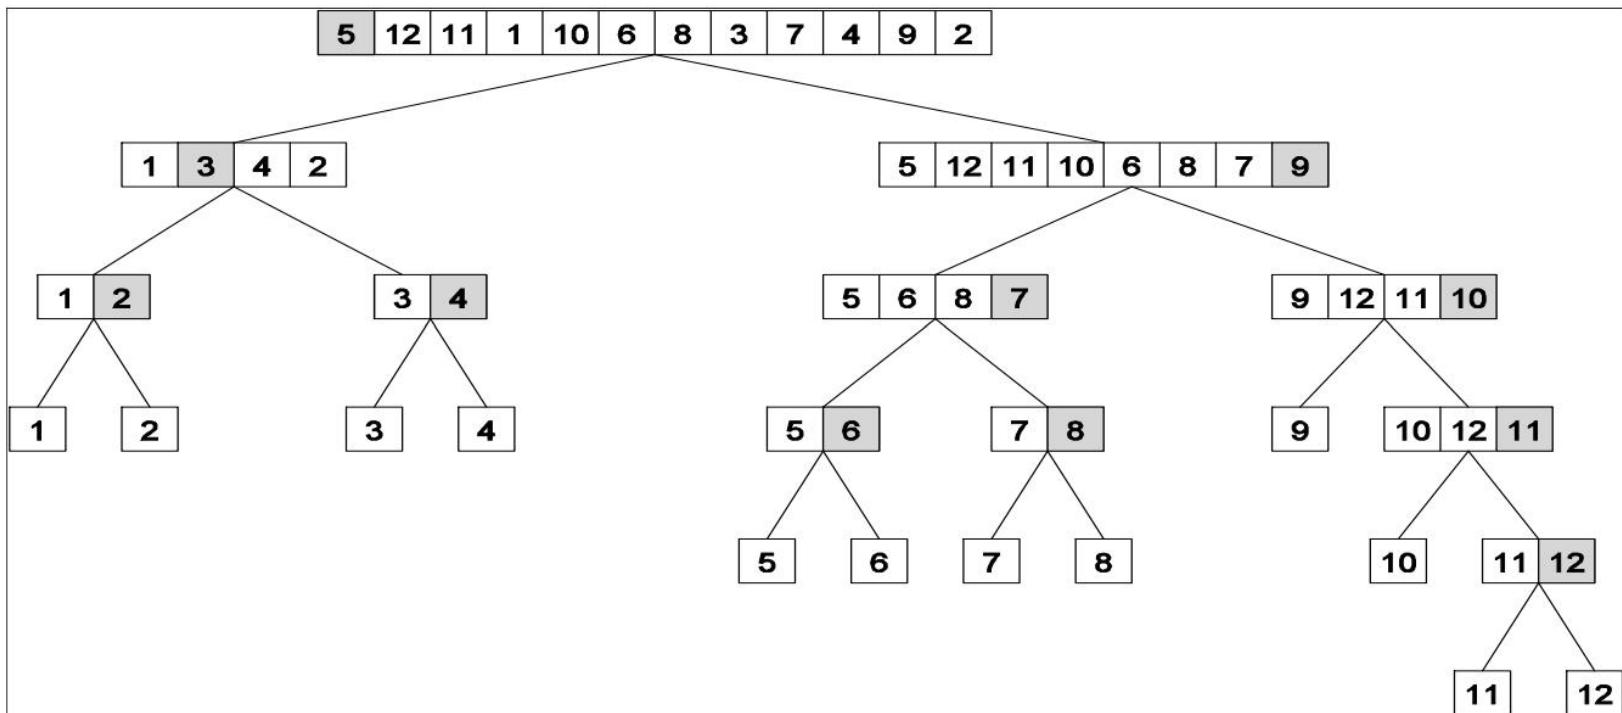
1. decompose a problem into a set of sub-problems
2. recursively decompose each sub-problem
3. stop decomposition when minimum desired granularity reached



Recursive Decomposition for Quicksort

Sort a vector v:

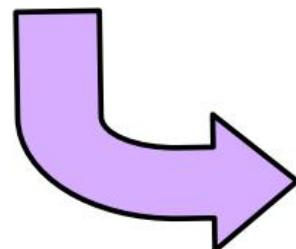
1. Select a pivot
2. Partition v around pivot into vleft and vright
3. In parallel, sort vleft and sort vright



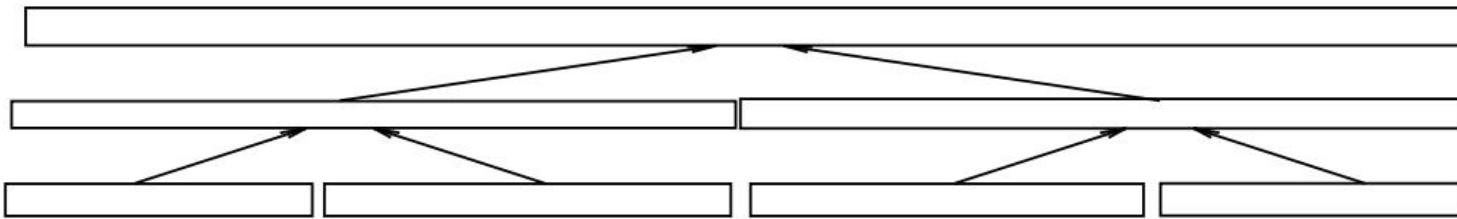
Source: COMP 422 at Rice Univ., Spring 2016.

Recursive Decomposition for Min

```
procedure SERIAL_MIN(A, n)
begin
  min = A[0];
  for i := 1 to n - 1 do
    if (A[i] < min) min := A[i];
  return min;
```



```
procedure RECURSIVE_MIN (A, n)
begin
  if (n = 1) then
    min := A[0];
  else
    lmin := spawn RECURSIVE_MIN(&A[0], n/2);
    rmin := spawn RECURSIVE_MIN(&A[n/2], n-n/2);
    if (lmin < rmin) then
      min := lmin;
    else
      min := rmin;
  return min;
```



Applicable to other associative operations, e.g. sum, AND ...

Exercise: Recursive Decomposition

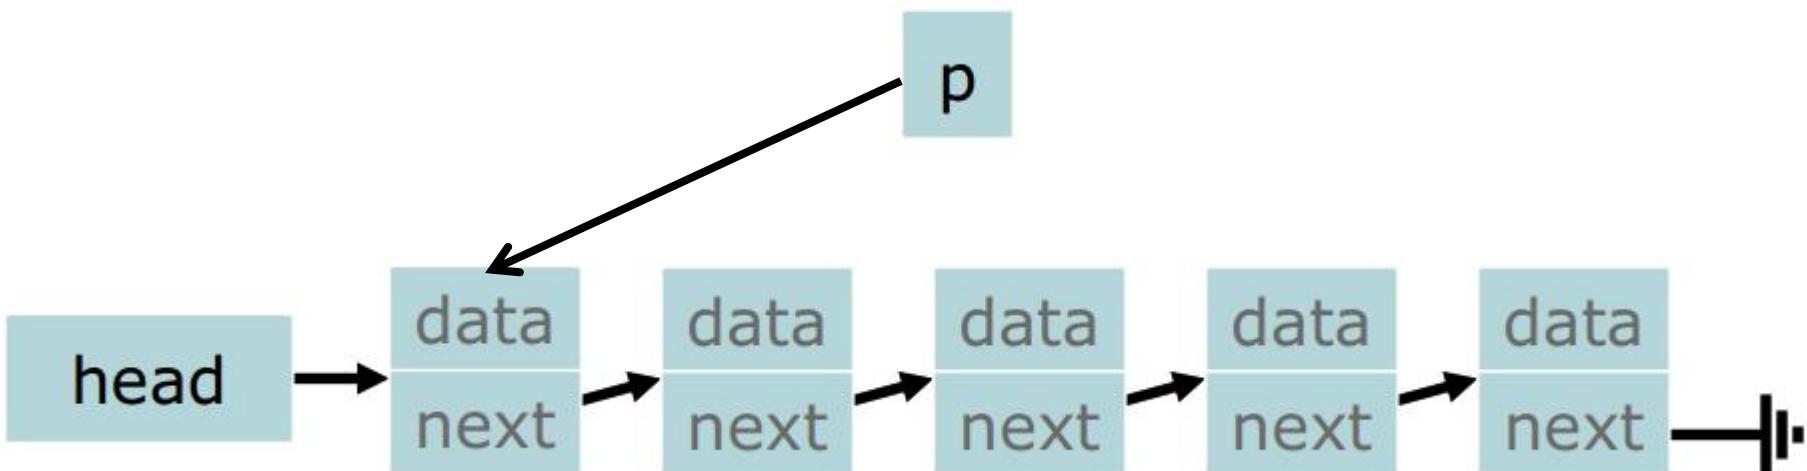
- ◆ **Recursive decomposition using OpenMP**
- ◆ **How to use OpenMP to do so?**
- ◆ **Tasks: allows parallelization of irregular problems**
 - **Unbounded loops**
 - **Recursive algorithms**
 - **Producer/consumer**

What are Tasks?

- ◆ **Tasks are independent units of work**
 - **Threads are assigned to perform the work of each task**
 - **Tasks may be deferred**
 - **Tasks may be executed immediately**
- ◆ **The runtime system decides which of the above**

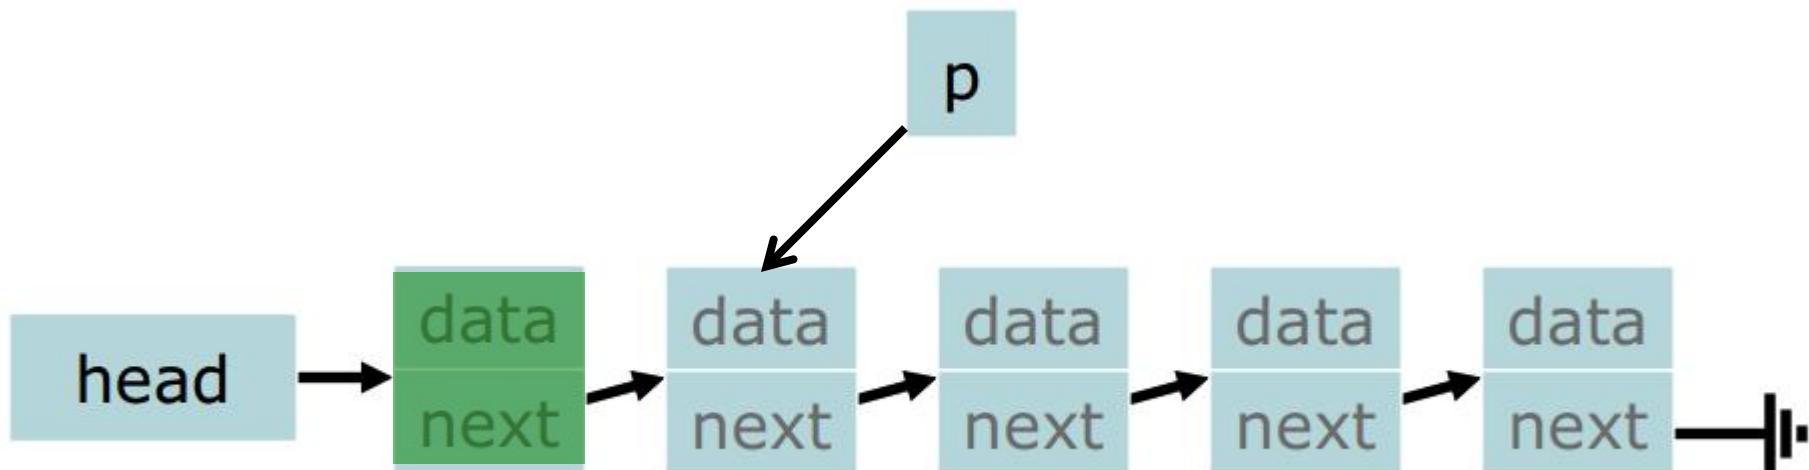
A Linked List Example

```
node* p = head;  
while (p) {  
    process(p);  
    p = p->next;  
}
```



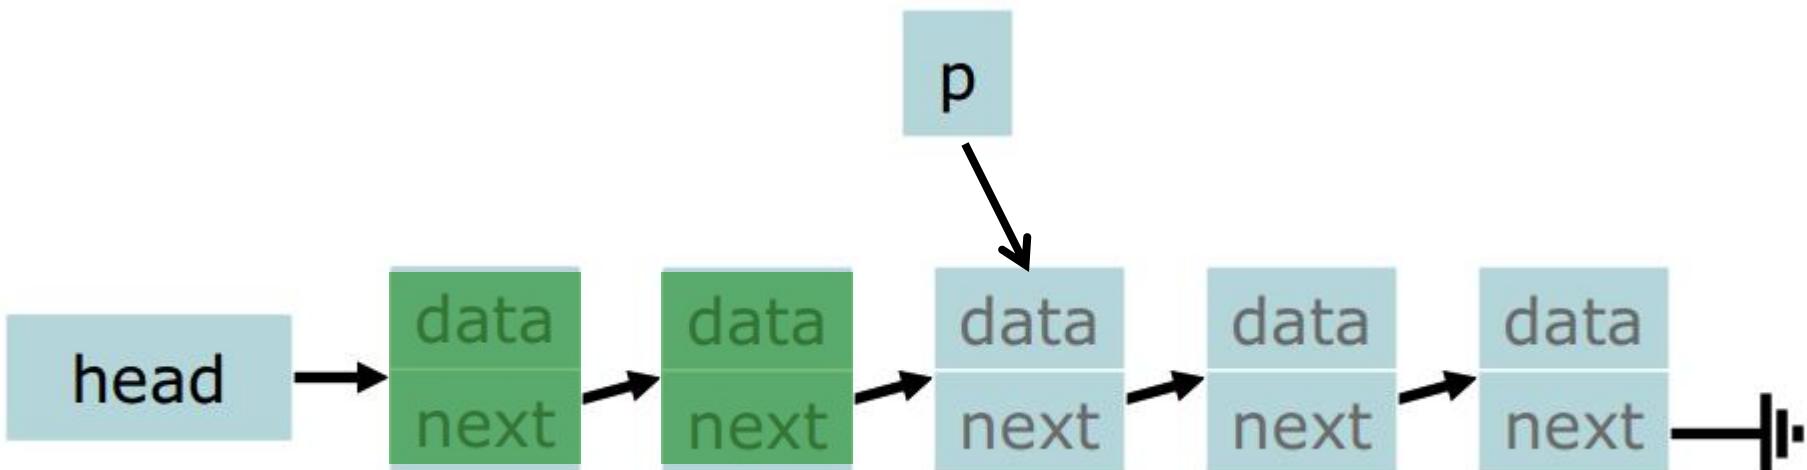
A Linked List Example

```
node* p = head;  
while (p) {  
    process(p);  
    p = p->next;  
}
```



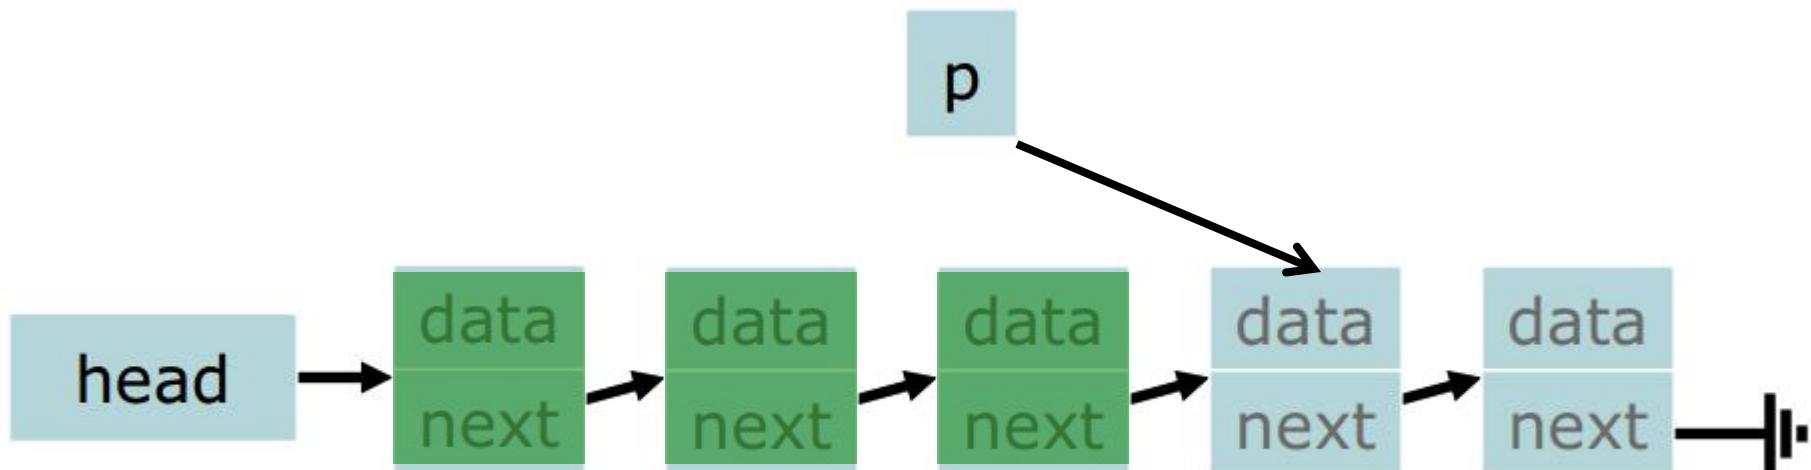
A Linked List Example

```
node* p = head;  
while (p) {  
    process(p);  
    p = p->next;  
}
```



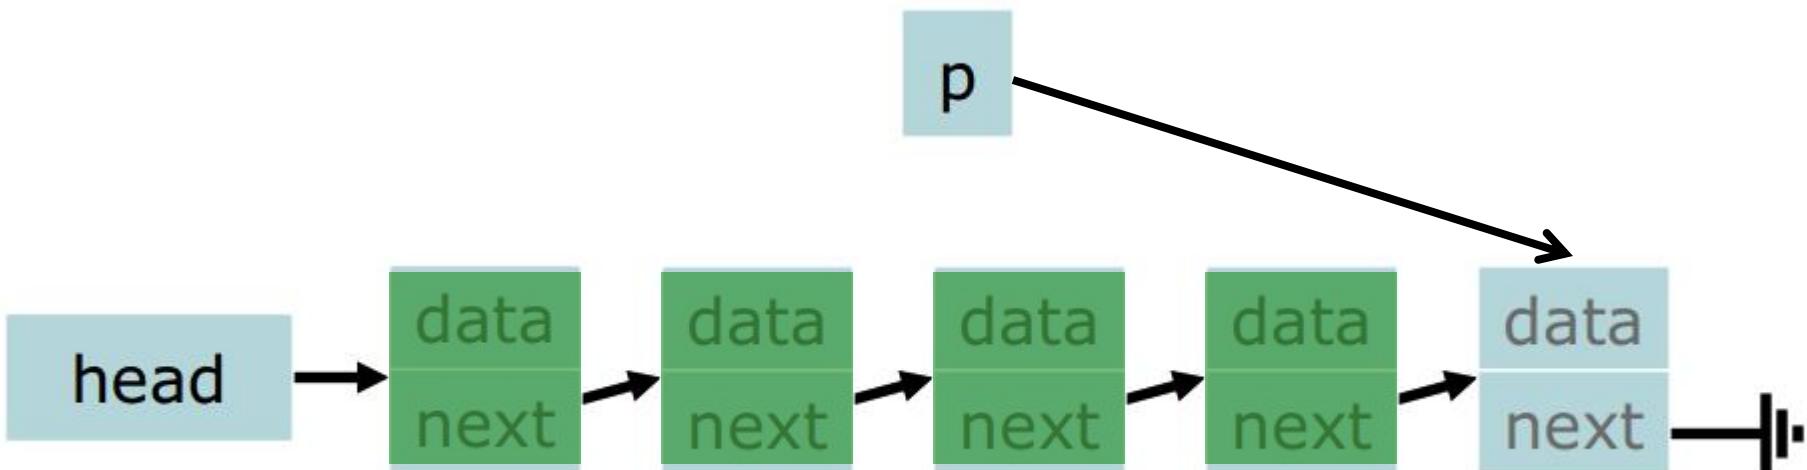
A Linked List Example

```
node* p = head;  
while (p) {  
    process(p);  
    p = p->next;  
}
```



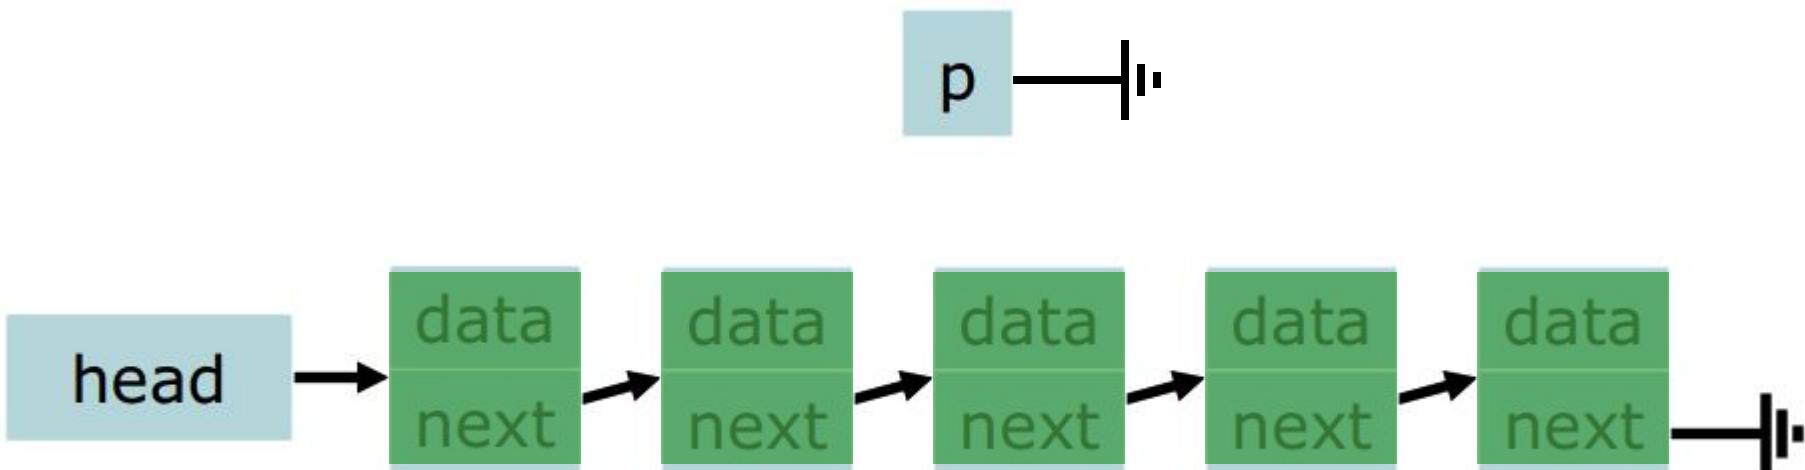
A Linked List Example

```
node* p = head;  
while (p) {  
    process(p);  
    p = p->next;  
}
```



A Linked List Example

```
node* p = head;  
while (p) {  
    process(p);  
    p = p->next;  
}
```

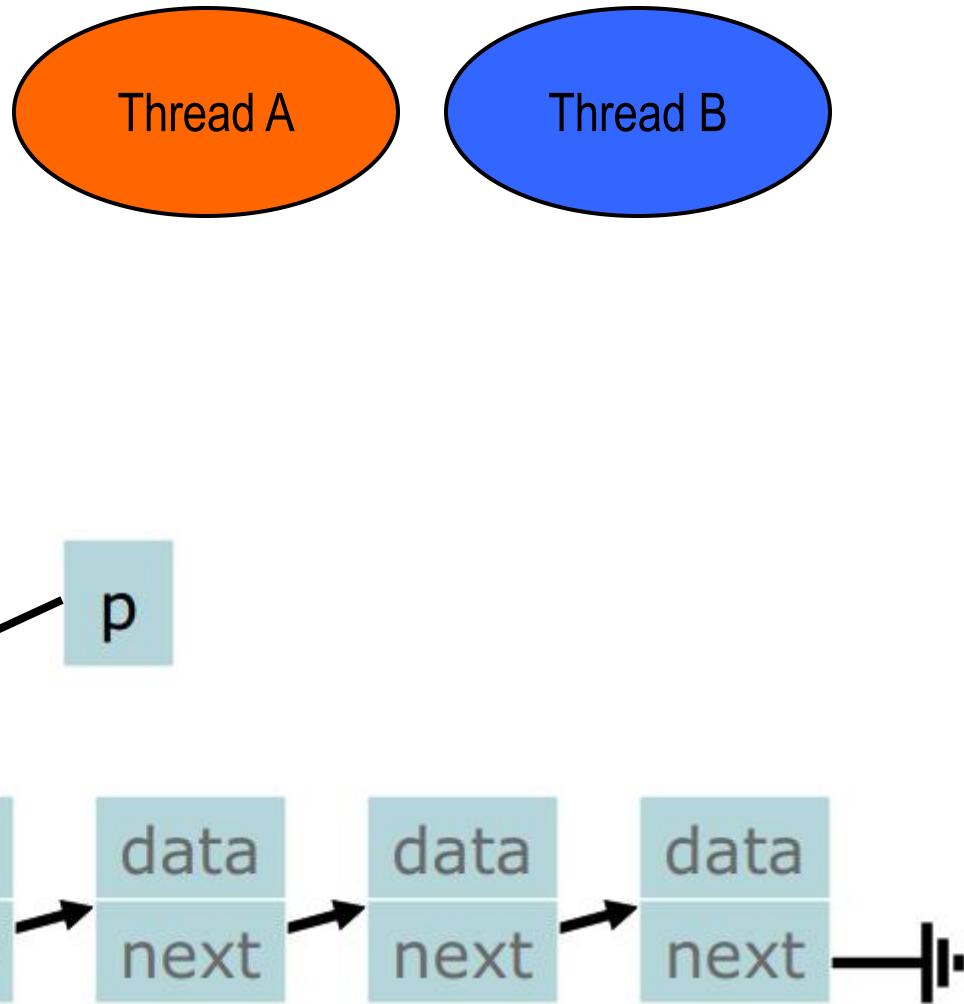


Task Construct – Explicit Task View

```
node* p = head;
#pragma omp parallel
// a team of threads is forked here
{
#pragma omp single
    // a single thread, T, executes the while loop
    while (p) {
#pragma omp task
        // each time T crosses the pragma, it generates a new task
        process(p); // each task runs in a thread
        p = p->next; // only T updates the p pointer
    } // all task complete at this implicit barrier of single
}
```

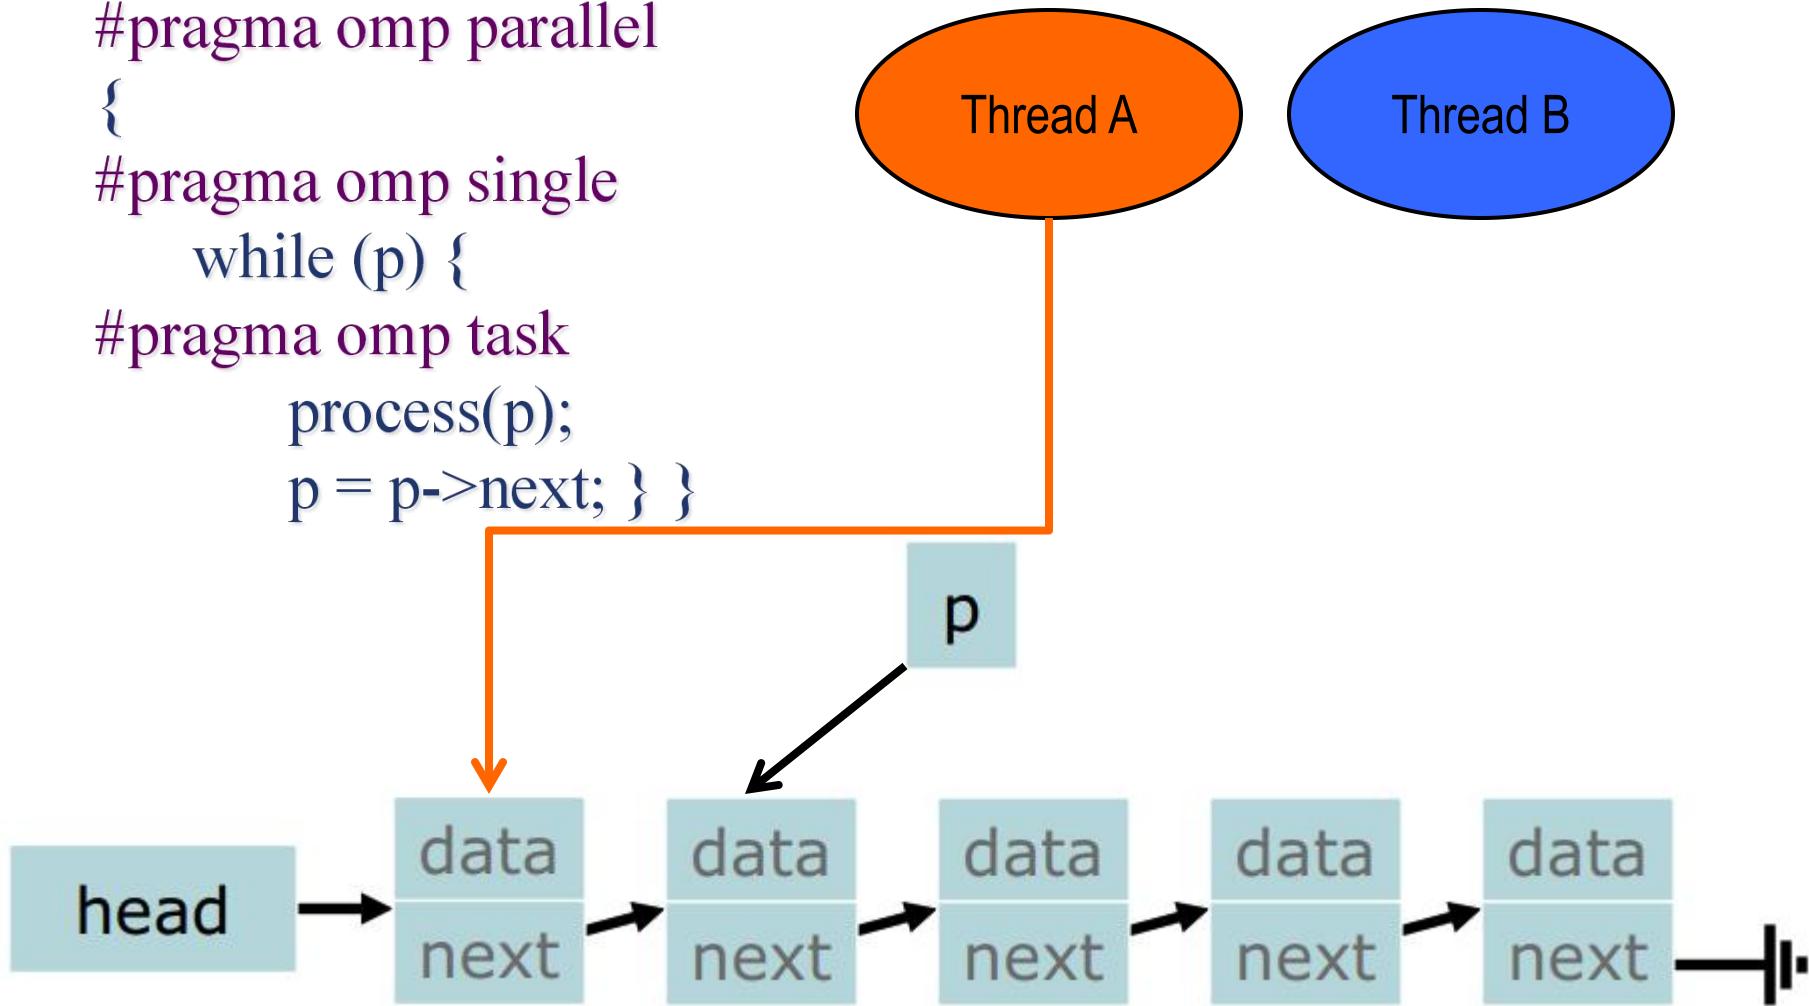
A Linked List Example

```
node* p = head;  
#pragma omp parallel  
{  
#pragma omp single  
    while (p) {  
#pragma omp task  
        process(p);  
        p = p->next; } }
```



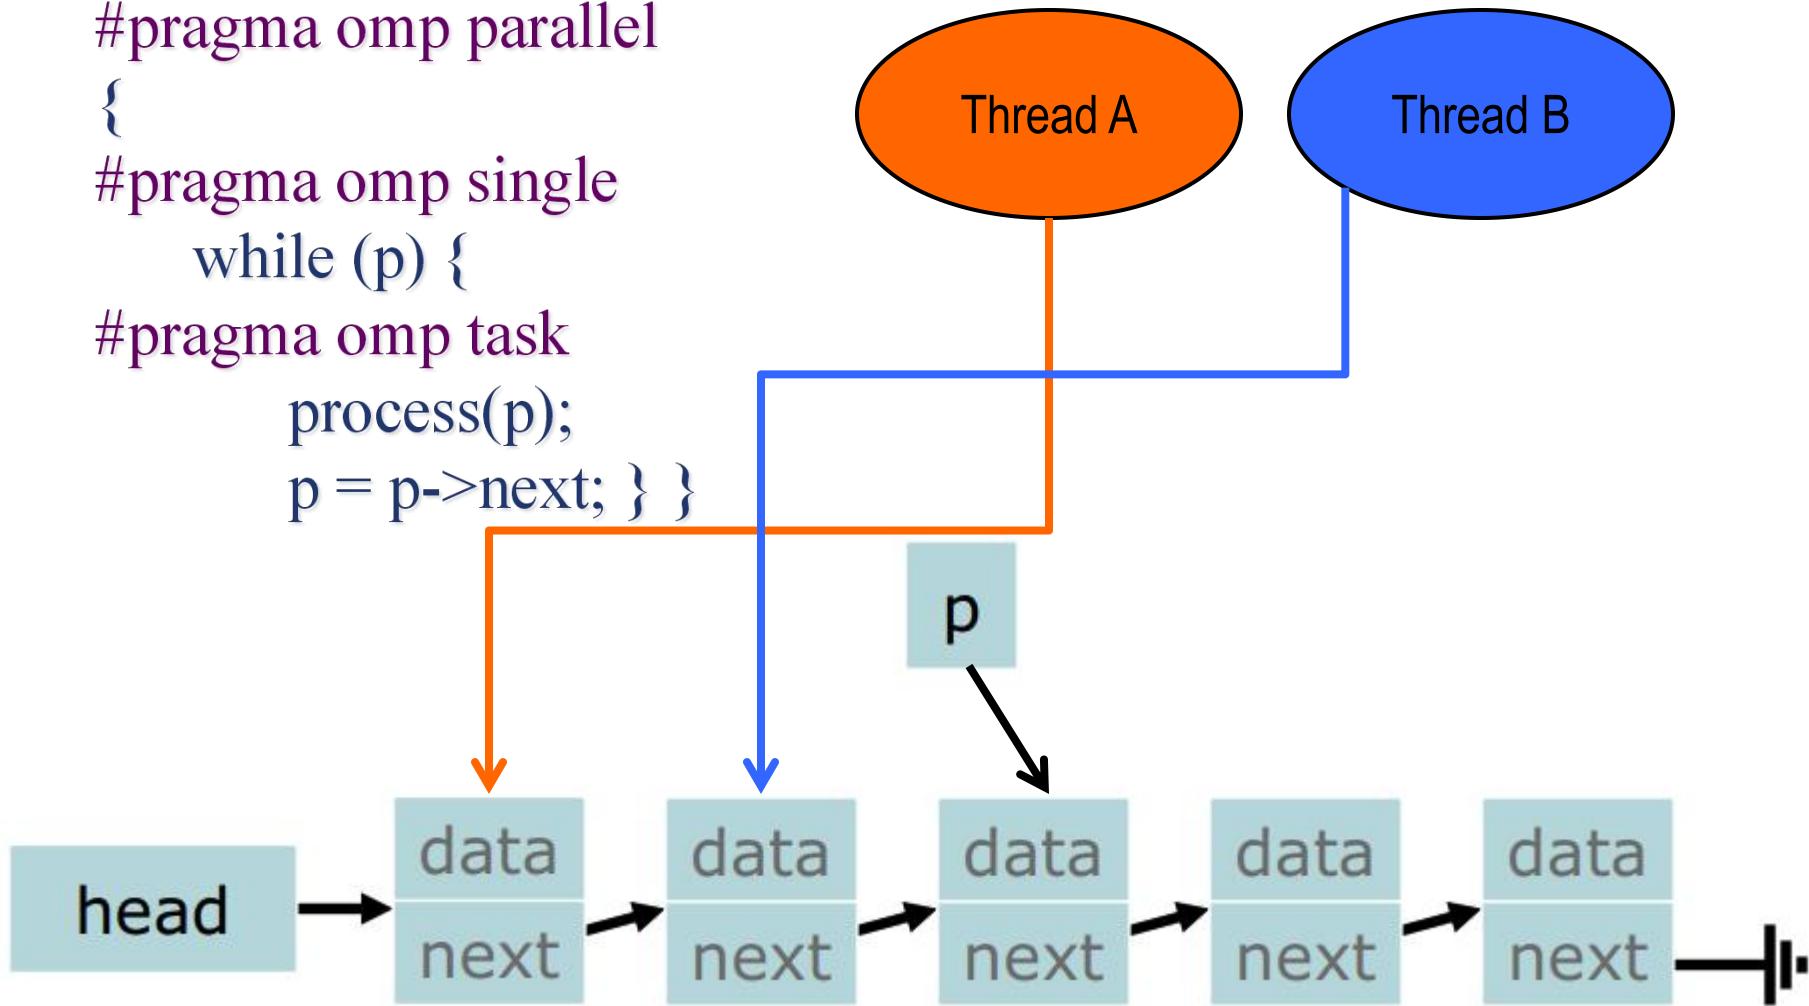
A Linked List Example

```
node* p = head;  
#pragma omp parallel  
{  
#pragma omp single  
    while (p) {  
#pragma omp task  
        process(p);  
        p = p->next; } }
```



A Linked List Example

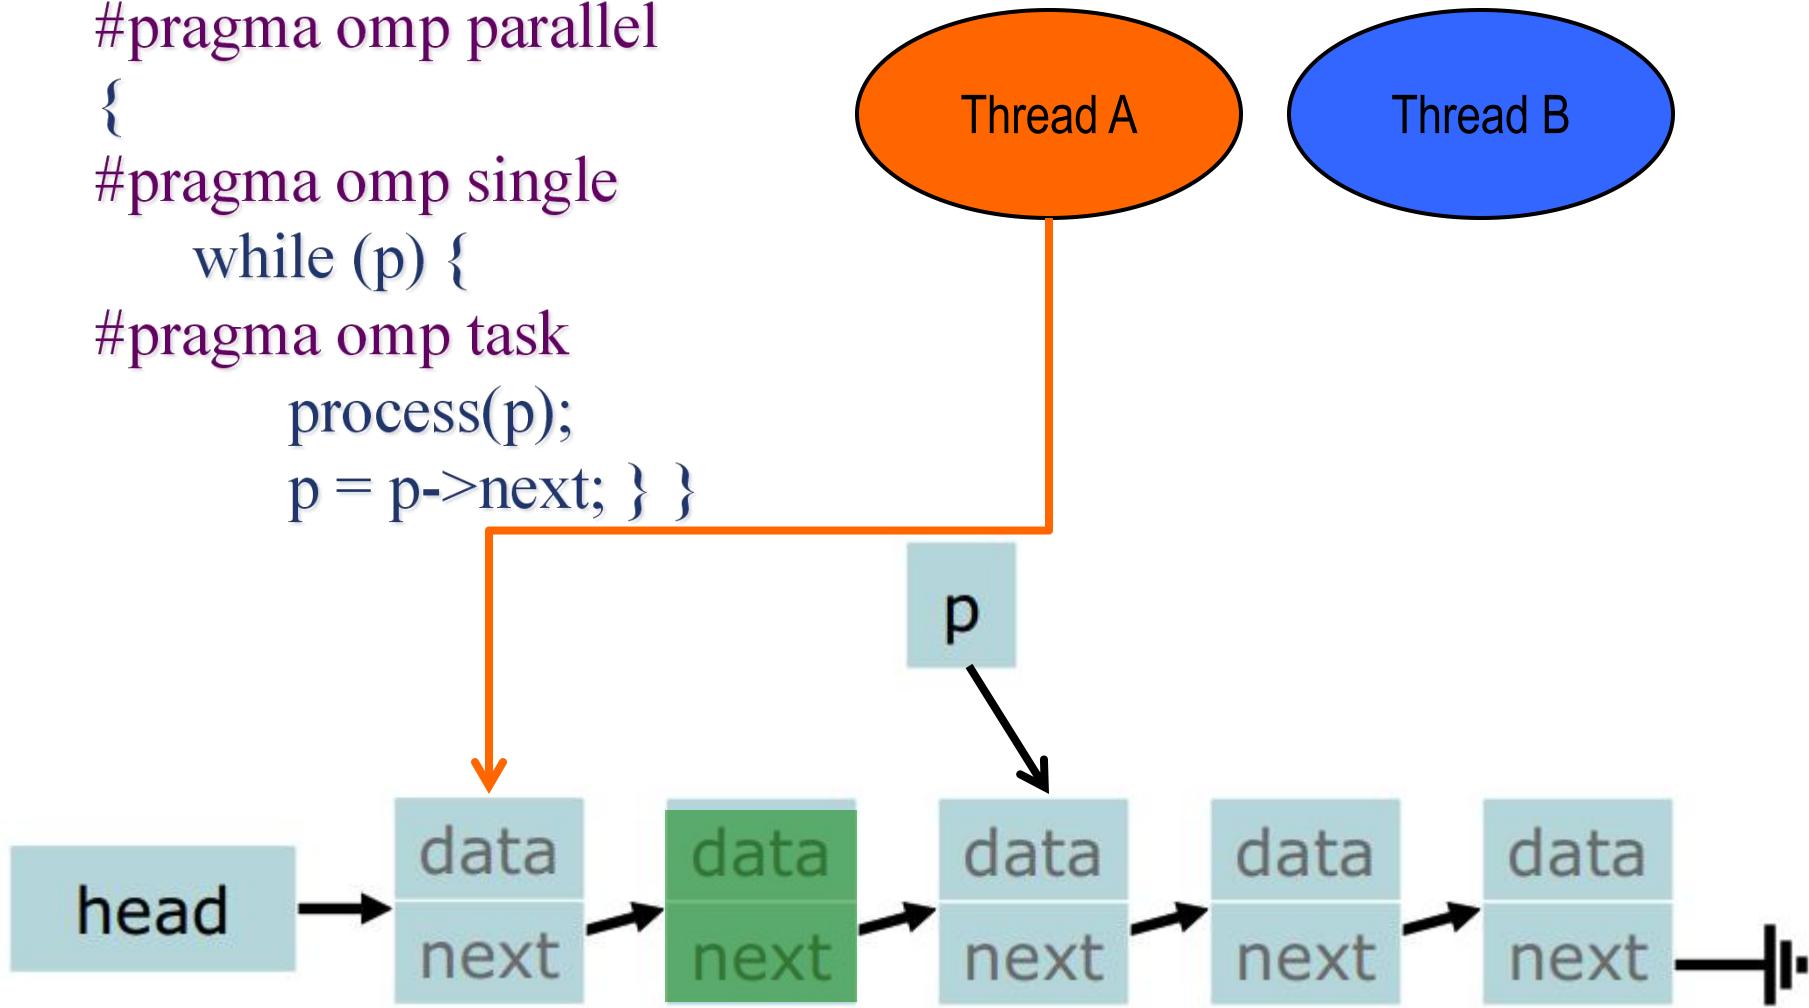
```
node* p = head;  
#pragma omp parallel  
{  
#pragma omp single  
    while (p) {  
#pragma omp task  
        process(p);  
        p = p->next; } }
```



Intel, "Introduction to Parallel Programming," 2009

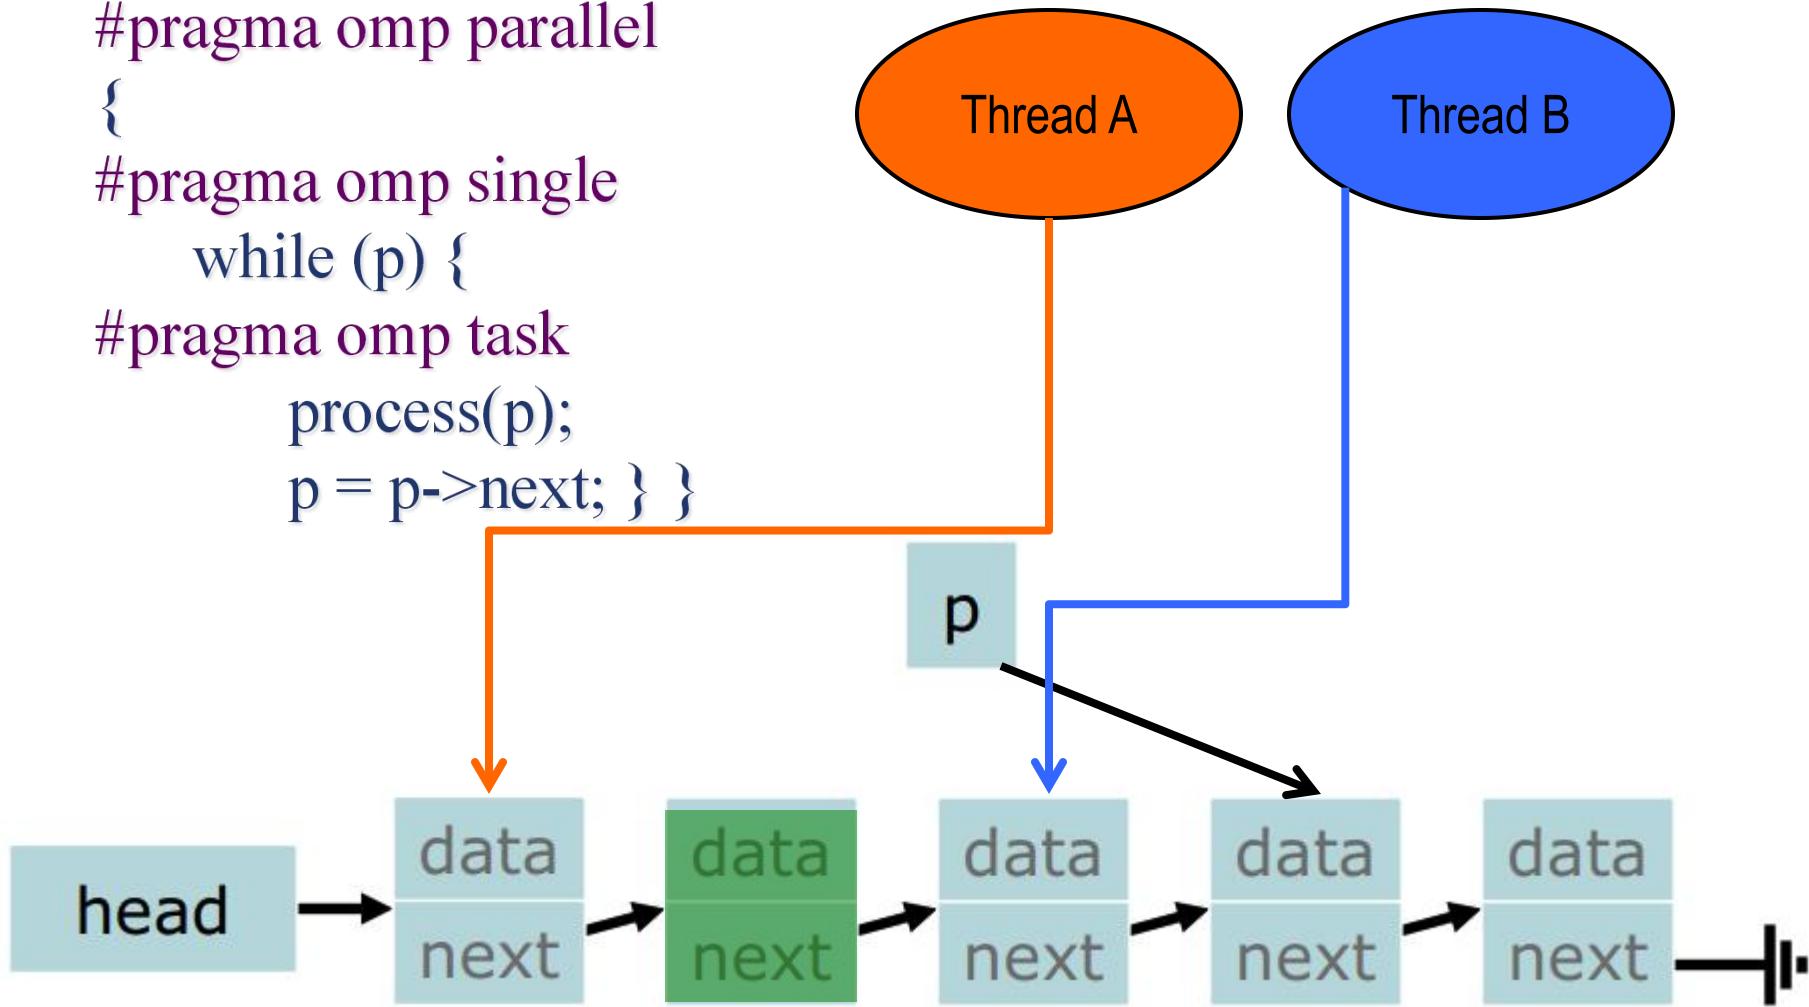
A Linked List Example

```
node* p = head;  
#pragma omp parallel  
{  
#pragma omp single  
    while (p) {  
#pragma omp task  
        process(p);  
        p = p->next; } }
```



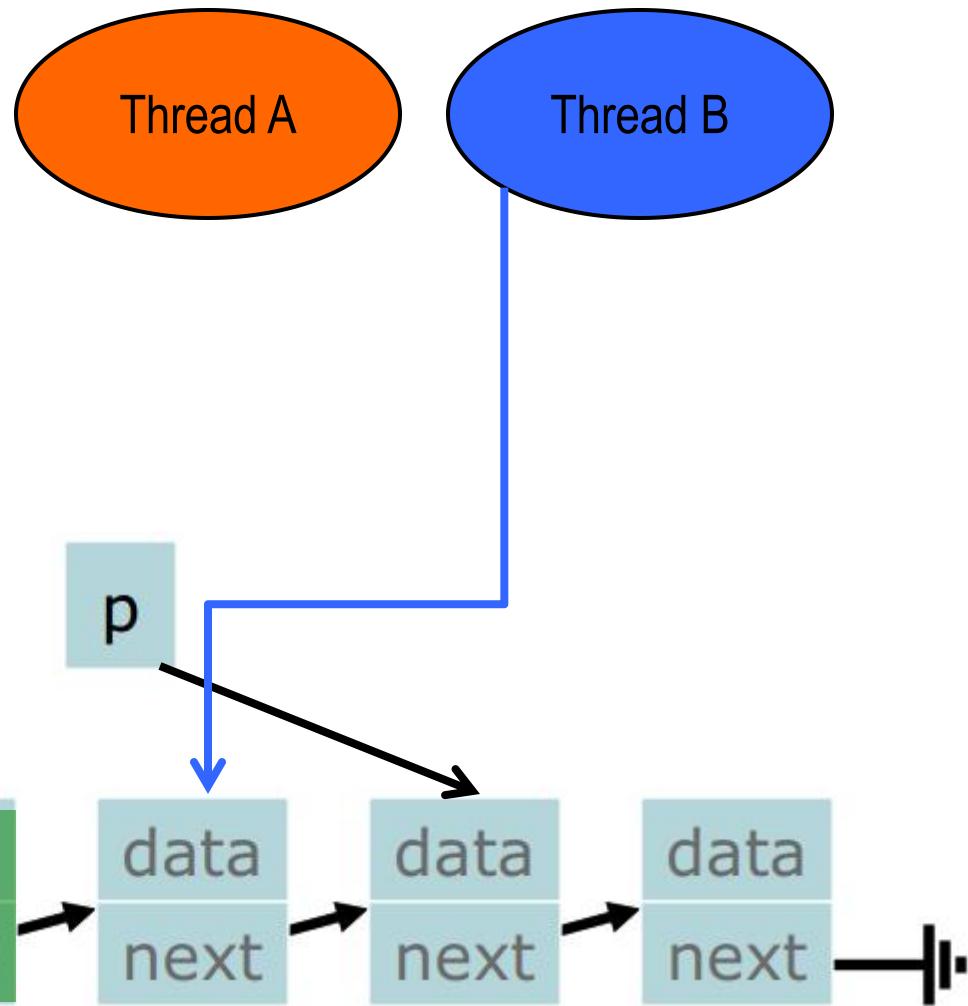
A Linked List Example

```
node* p = head;  
#pragma omp parallel  
{  
#pragma omp single  
    while (p) {  
#pragma omp task  
        process(p);  
        p = p->next; } }
```



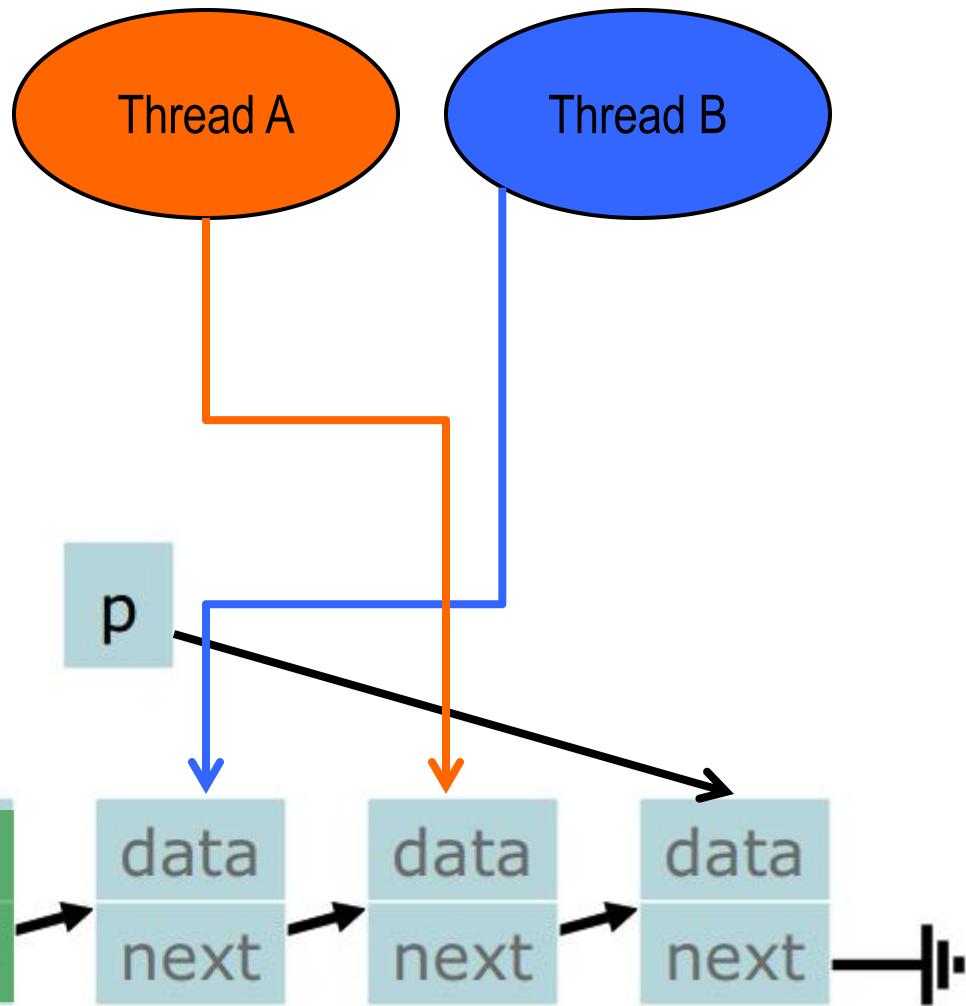
A Linked List Example

```
node* p = head;  
#pragma omp parallel  
{  
#pragma omp single  
    while (p) {  
#pragma omp task  
        process(p);  
        p = p->next; } }
```



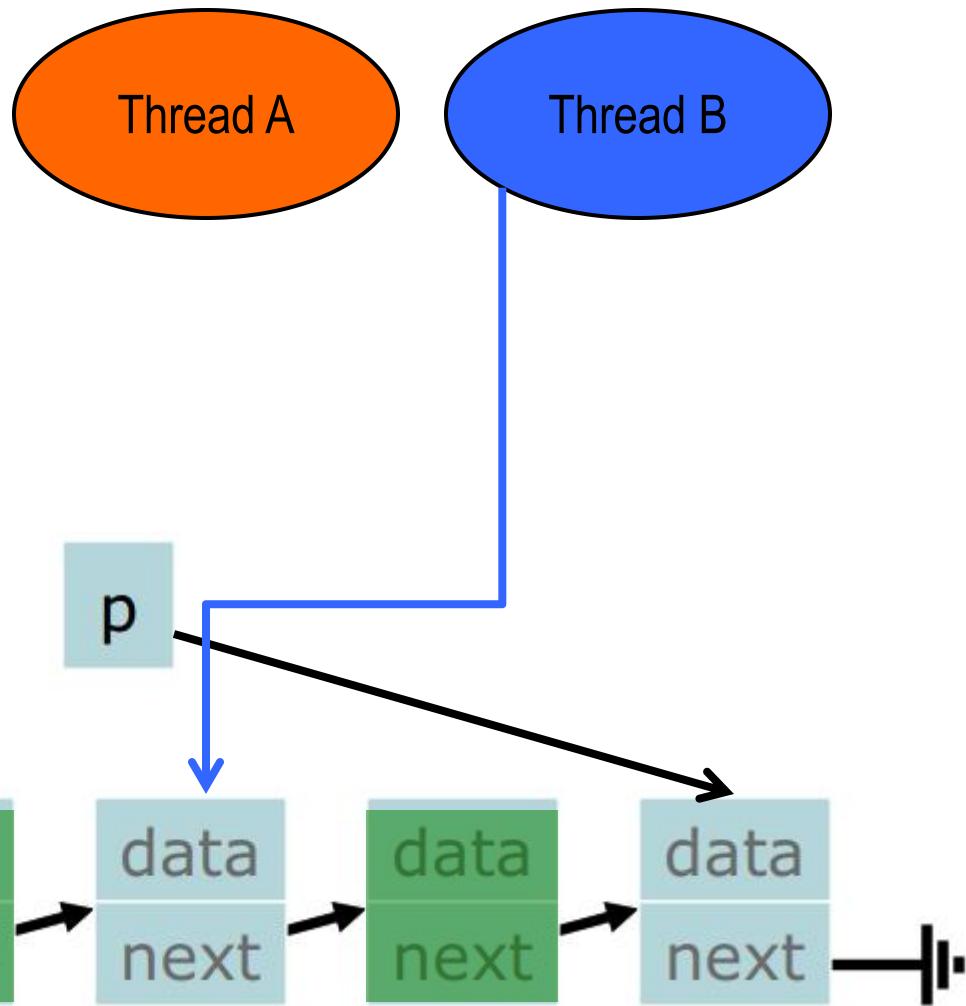
A Linked List Example

```
node* p = head;  
#pragma omp parallel  
{  
#pragma omp single  
    while (p) {  
#pragma omp task  
        process(p);  
        p = p->next; } }
```



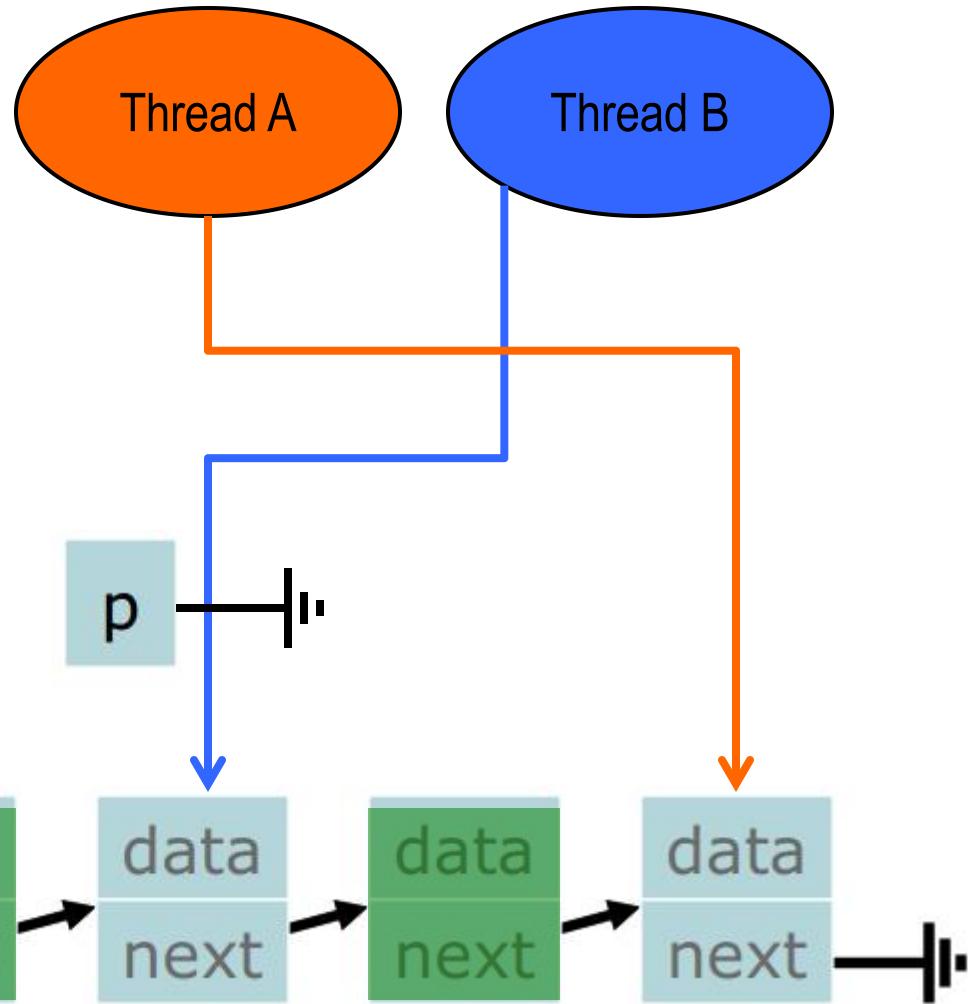
A Linked List Example

```
node* p = head;  
#pragma omp parallel  
{  
#pragma omp single  
    while (p) {  
#pragma omp task  
        process(p);  
        p = p->next; } }
```



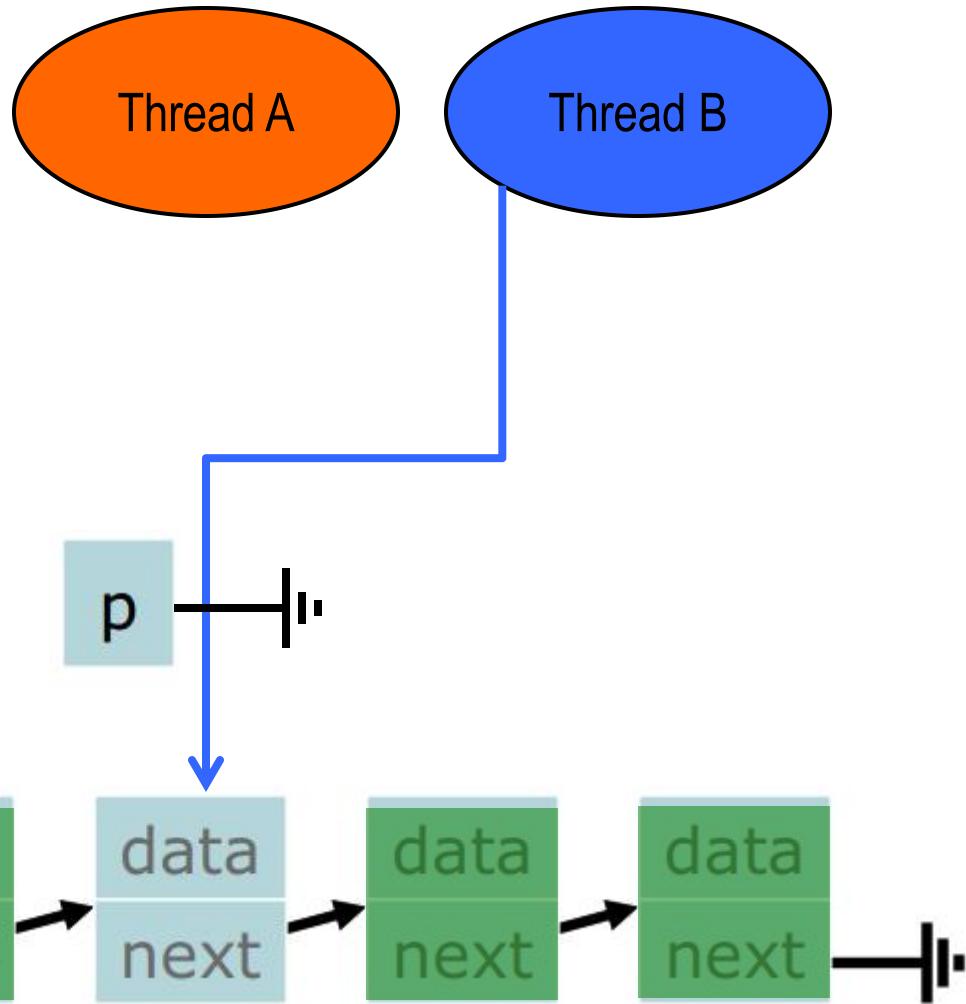
A Linked List Example

```
node* p = head;  
#pragma omp parallel  
{  
#pragma omp single  
    while (p) {  
#pragma omp task  
        process(p);  
        p = p->next; } }
```



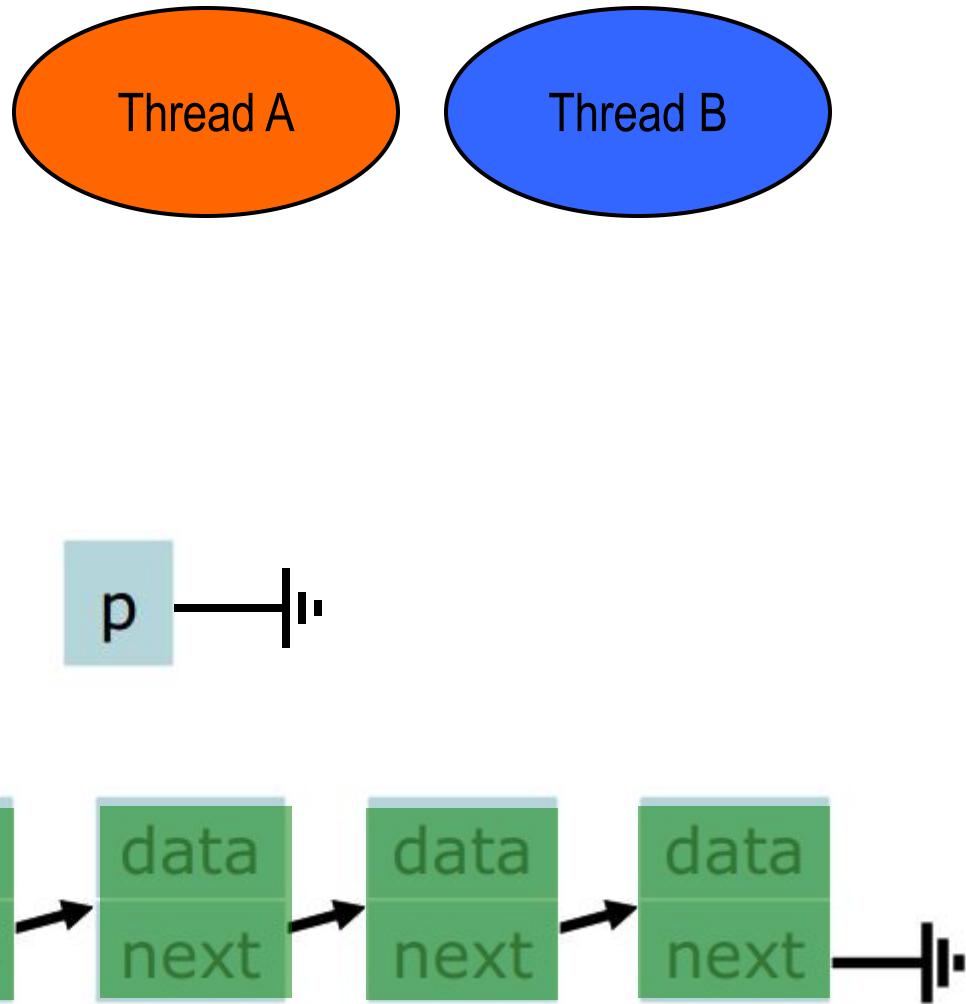
A Linked List Example

```
node* p = head;  
#pragma omp parallel  
{  
#pragma omp single  
    while (p) {  
#pragma omp task  
        process(p);  
        p = p->next; } }
```



A Linked List Example

```
node* p = head;  
#pragma omp parallel  
{  
#pragma omp single  
    while (p) {  
#pragma omp task  
        process(p);  
        p = p->next; } }
```



When are Tasks Guaranteed to be Complete?

◆ Tasks are guaranteed to be complete:

- At thread or task barriers
- At the directive: **#pragma omp barrier**
 - all the threads in the team will be synchronized
- At the directive: **#pragma omp taskwait**
 - only the child threads generated since the beginning of the current task will be synchronized

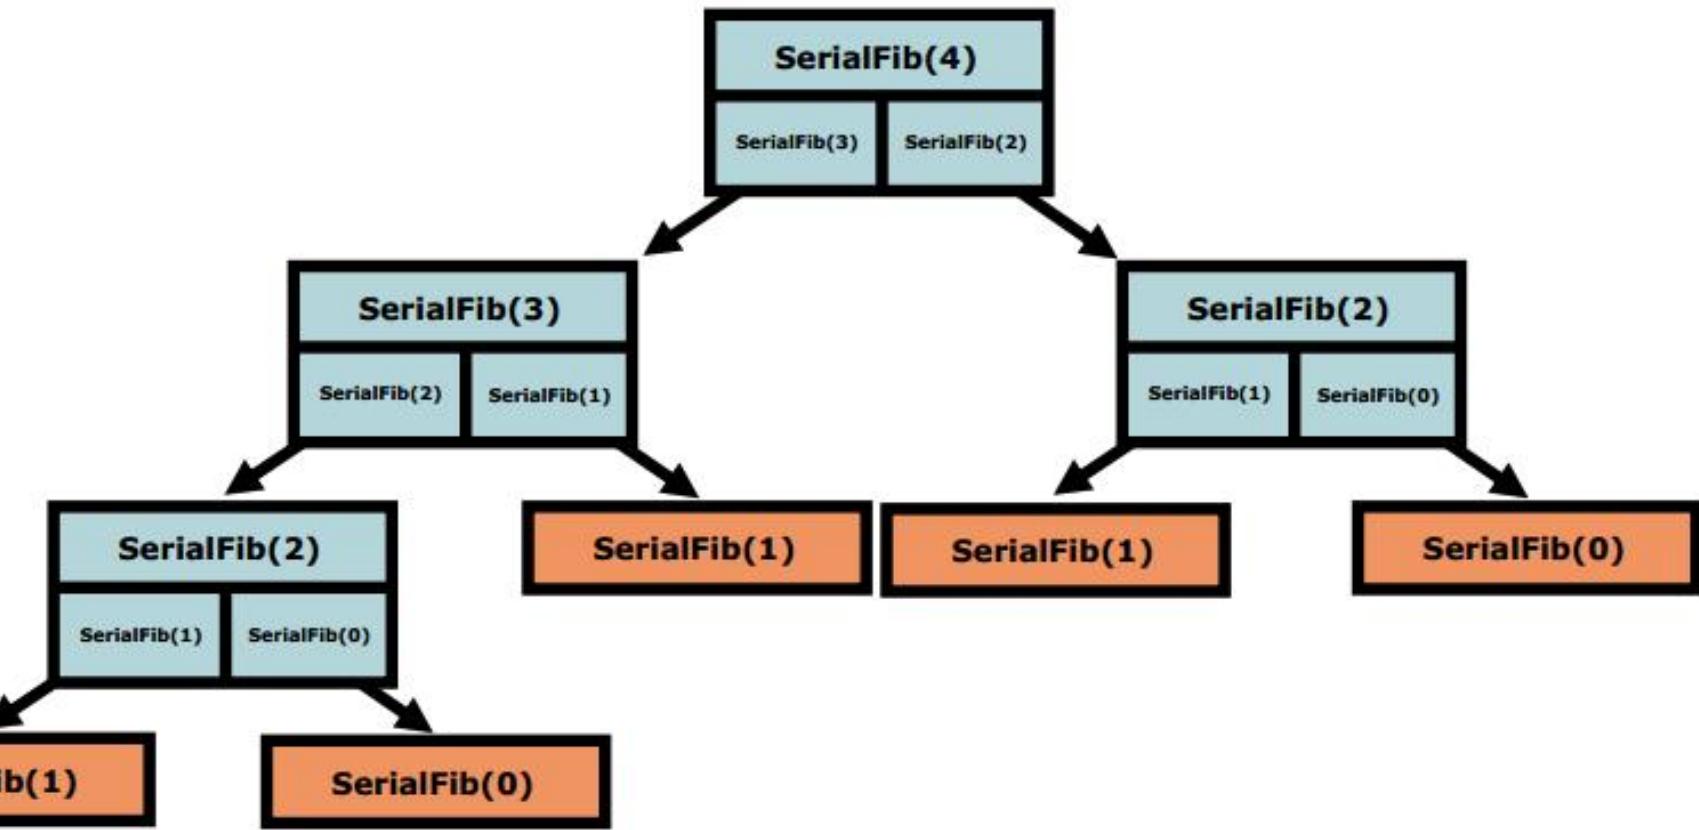
Example: Naïve Fibonacci Calculation

- ◆ Recursion typically used to calculate Fibonacci number
- ◆ Widely used as toy benchmark
 - Easy to code
 - Has unbalanced task graph

```
long SerialFib(long n) {  
    if (n < 2) return n;  
    else return SerialFib(n-1) + SerialFib(n-2);  
}
```

Example: Naïve Fibonacci Calculation

- ◆ We can envision Fibonacci computation as a task graph



Intel, "Introduction to Parallel Programming," 2009

Fibonacci – Task Spawning Solution

```
long ParallelFib(long n) {  
    long sum;  
#pragma omp parallel  
    {  
#pragma omp single  
        FibTask(n, &sum);  
    }  
    return sum;  
}
```

```
void FibTask(  
    long n, long* sum) {  
    if (n < CutOff)  
        *sum = SerialFib(n);  
    else {  
        long x, y;  
#pragma omp task shared(x)  
        FibTask(n-1, &x);  
#pragma omp task shared(y)  
        FibTask(n-2, &y);  
#pragma omp taskwait  
        *sum = x + y;  
    }  
}
```

Exploratory Decomposition

- ◆ Exploration (search) of a state space of solutions
 - problem decomposition reflects shape of execution
- ◆ Examples
 - discrete optimization
 - 0/1 integer programming
 - theorem proving
 - game playing

Exploratory Decomposition Example

Solving a 15 puzzle

- ◆ Sequence of three moves from state (a) to final state (d)

1	2	3	4
5	6	7	8
9	10	15	11
13	14	12	

(a)

1	2	3	4
5	6	7	8
9	10	11	15
13	14	12	

(b)

1	2	3	4
5	6	7	8
9	10	11	15
13	14	12	

(c)

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	

(d)

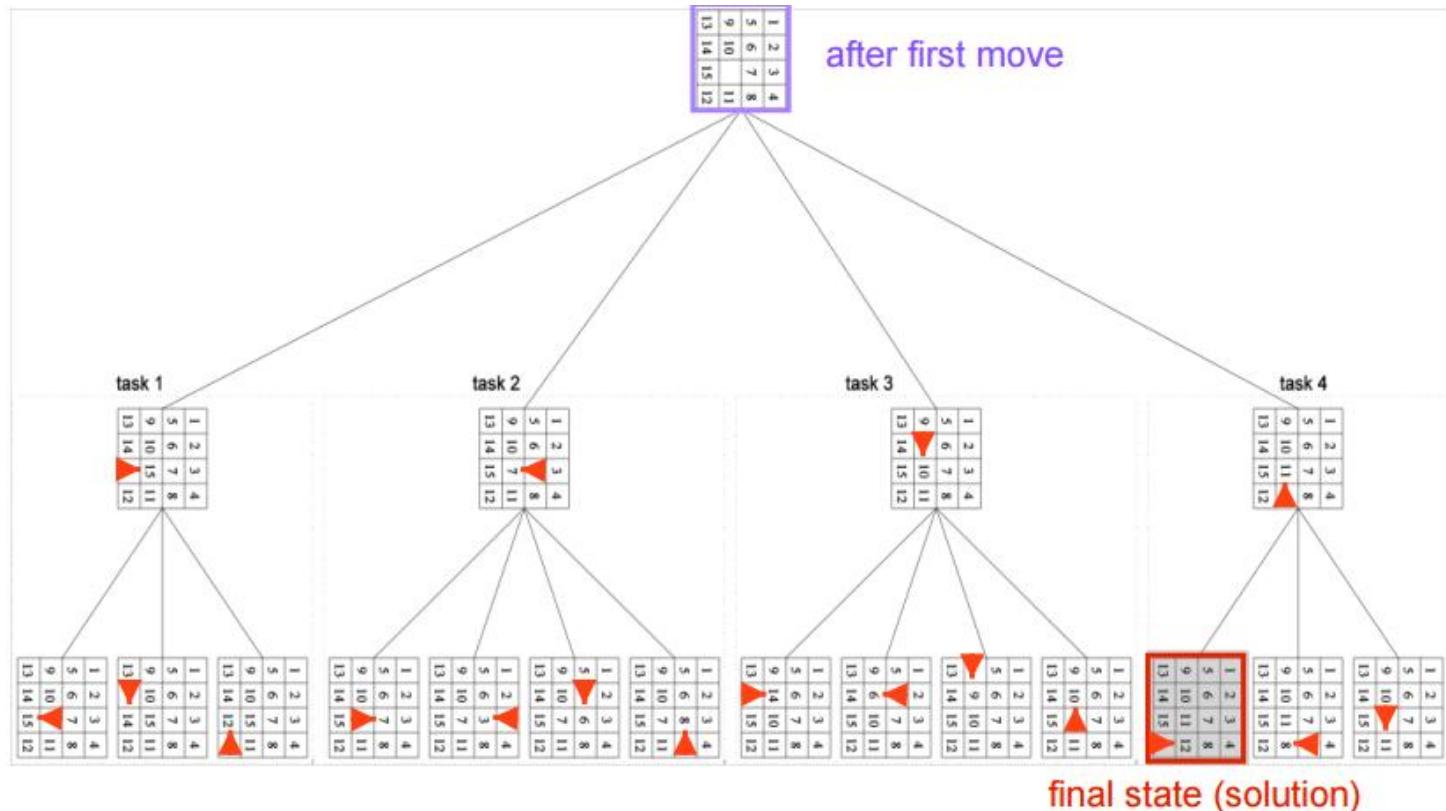
- ◆ From an arbitrary state, must search for a solution

Exploratory Decomposition: Example

Solving a 15 puzzle

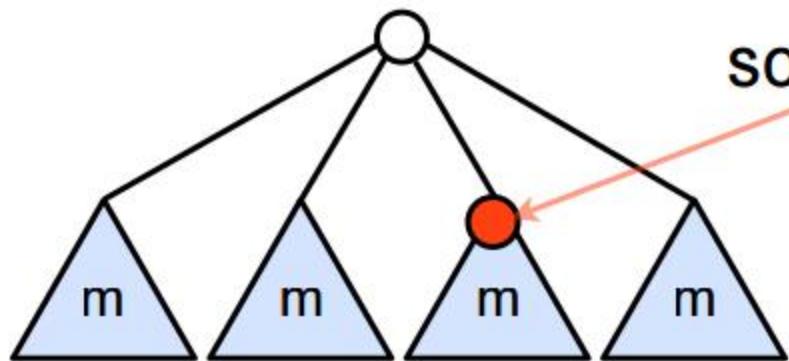
Search

- generate successor states of the current state
- explore each as an independent task

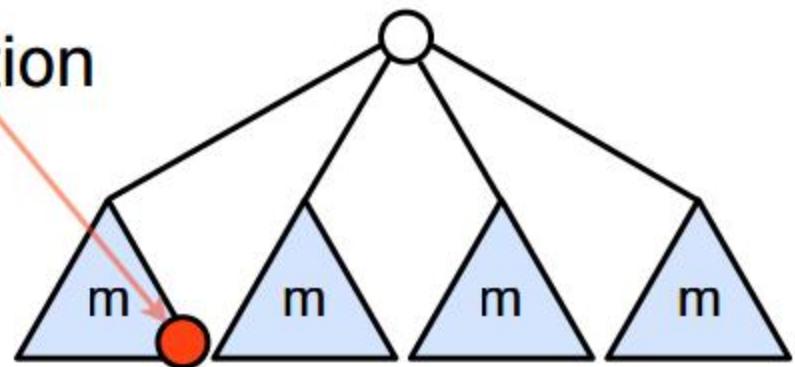


Exploratory Decomposition Speedup

- ◆ Parallel formulation may perform a different amount of work



total serial work = $2m + 1$
total parallel work = 4



total serial work = m
total parallel work = 4m

- ◆ Can cause super- or sub-linear speedup

Exploratory Decomposition

◆ Differences from data/recursive decomposition

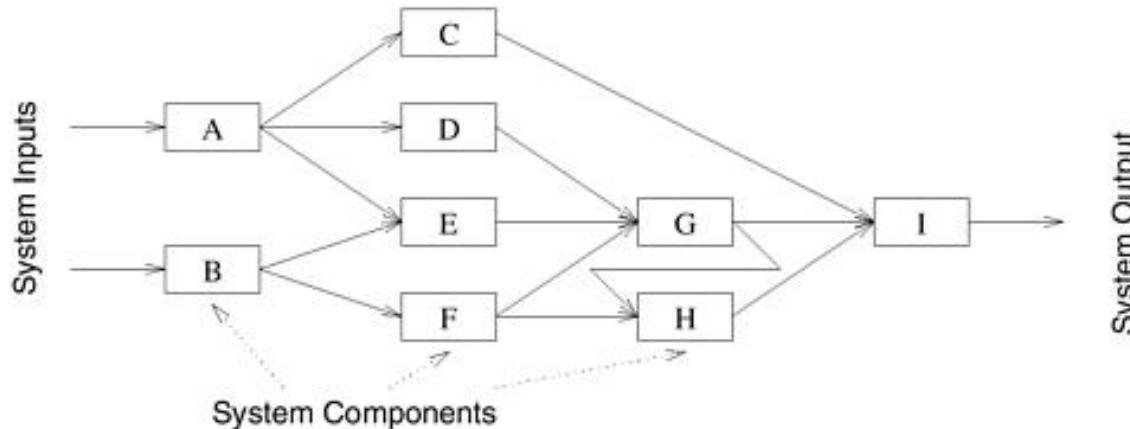
- Appear similar to data/recursive decomposition, if the search space is viewed as the data to get partitioned
- The tasks in data/recursive decomposition perform useful computations
- Unfinished tasks in exploratory decomposition can be terminated as soon as an overall solution is found

Speculative Decomposition

- ◆ Dependencies between tasks are not always known a-priori
 - makes it impossible to identify independent tasks
- ◆ Conservative approach
 - identify independent tasks only when no dependencies left
- ◆ Optimistic (speculative) approach
 - schedule tasks even when they may potentially be erroneous
- ◆ Drawbacks for each
 - conservative approaches
 - may yield little concurrency
 - optimistic approaches
 - may require a roll-back mechanism if a dependence is encountered

Speculative Decomposition

- ◆ Analogy: if/switch statements in instruction-level parallelism (ILP)
- ◆ A simple network for discrete event simulation



- ◆ Speculative vs exploratory
 - speculative: input is unknown; may perform more or the same amount of work
 - exploratory: output is unknown; may perform more, less, or the same amount of work

Speculative Decomposition In Practice

Discrete event simulation

- ◆ **Data structure: centralized time-ordered event list**
- ◆ **Simulation**
 - extract next event in time order
 - process the event
 - if required, insert new events into the event list
- ◆ **Optimistic event scheduling**
 - assume outcomes of all prior events
 - speculatively process next event
 - if assumption is incorrect, roll back its effects and continue

Time Warp

David Jefferson. “Virtual Time,”
ACM TOPLAS, 7(3):404-425, July 1985

Summary

◆ Decomposition techniques

- data decomposition**
- recursive decomposition**
- exploratory decomposition**
- speculative decomposition**