



# Introduction to Parallel and Distributed Computing

## Selected Parallel Algorithms Dynamic Programming

Lecture 12, Spring 2024

Instructor: 罗国杰

[gluo@pku.edu.cn](mailto:gluo@pku.edu.cn)

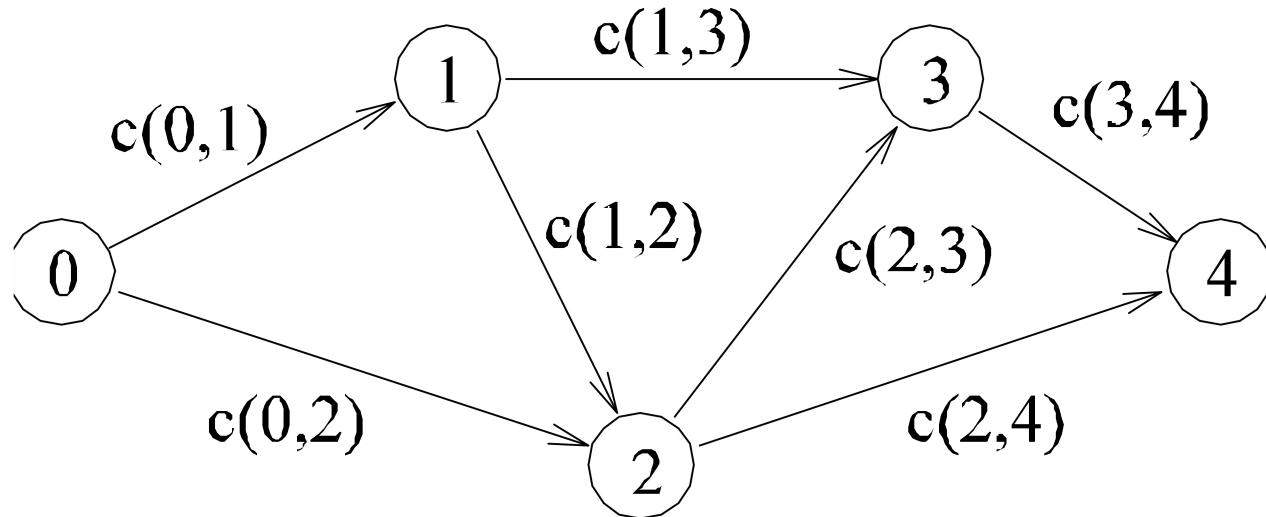
# ***Overview of Serial Dynamic Programming***

---

- ♦ Dynamic programming (DP) is used to solve a wide variety of discrete optimization problems such as scheduling, string-editing, packaging, and inventory management.
- ♦ Break problems into subproblems and combine their solutions into solutions to larger problems.
- ♦ In contrast to divide-and-conquer, there may be relationships across subproblems.

# **Dynamic Programming: Example**

- ◆ An acyclic graph for which the shortest path between nodes 0 and 4 is to be computed.



$$f(4) = \min\{f(3) + c(3,4), f(2) + c(2,4)\}.$$

# Dynamic Programming: Example

- ♦ Consider the problem of finding a shortest path between a pair of vertices in an acyclic graph.
- ♦ An edge connecting node  $i$  to node  $j$  has cost  $c(i,j)$ .
- ♦ The graph contains  $n$  nodes numbered  $0,1,\dots,n-1$ , and has an edge from node  $i$  to node  $j$  only if  $i < j$ . Node  $0$  is source and node  $n-1$  is the destination.
- ♦ Let  $f(x)$  be the cost of the shortest path from node  $0$  to node  $x$ .

$$f(x) = \begin{cases} 0 & x = 0 \\ \min_{0 \leq j < x} \{ f(j) + c(j,x) \} & 1 \leq x \leq n - 1 \end{cases}$$

# Dynamic Programming

- ♦ The solution to a DP problem is typically expressed as a minimum (or maximum) of possible alternate solutions.
- ♦ If  $r$  represents the cost of a solution composed of subproblems  $x_1, x_2, \dots, x_l$ , then  $r$  can be written as

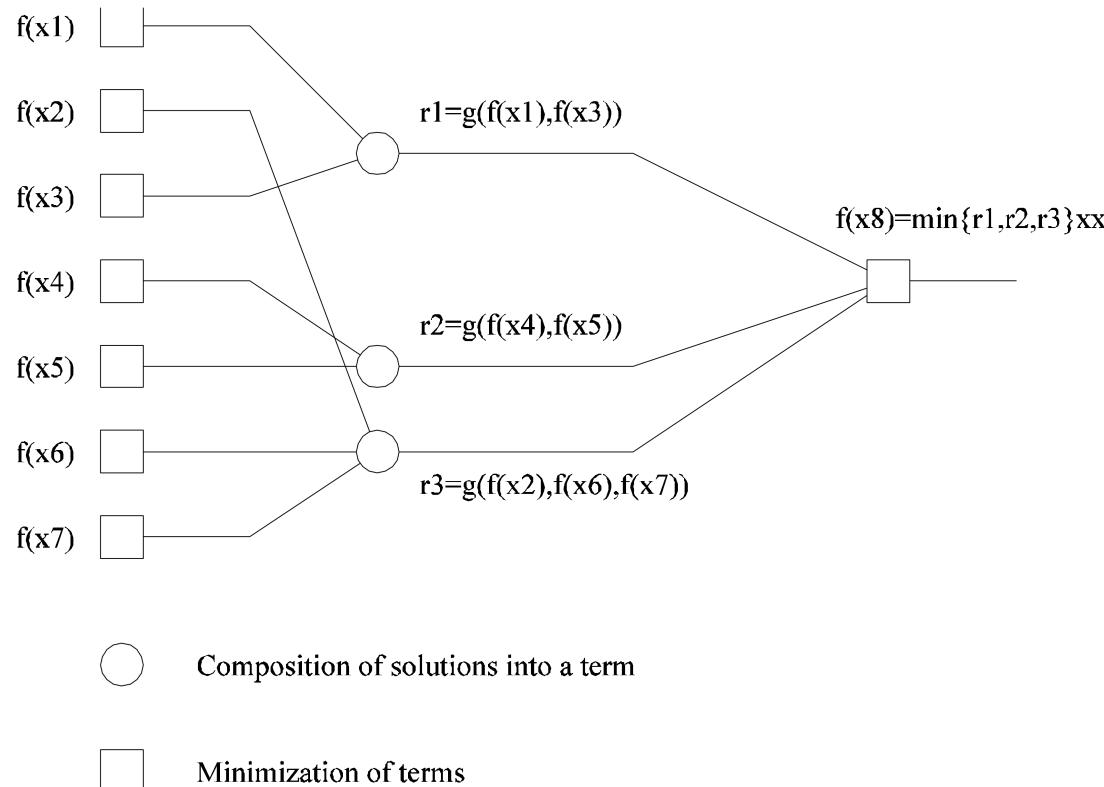
$$r = g(f(x_1), f(x_2), \dots, f(x_l)).$$

Here,  $g$  is the composition function.

- ♦ If the optimal solution to each problem is determined by composing optimal solutions to the subproblems and selecting the minimum (or maximum), the formulation is said to be a DP formulation.

# Dynamic Programming: Example

- ♦ The computation and composition of subproblem solutions to solve problem  $f(x_8)$ .



# **DP: Monadic vs. Polyadic**

---

- ♦ The recursive DP equation is also called the functional equation or optimization equation.
- ♦ In the equation for the shortest path problem the composition function is  $f(j) + c(j,x)$ . This contains a single recursive term ( $f(j)$ ). Such a formulation is called **monadic**.
- ♦ If the RHS has multiple recursive terms, the DP formulation is called **polyadic**.

## ***DP: Serial vs. Non-serial***

---

- ♦ The dependencies between subproblems can be expressed as a graph.
- ♦ If the graph can be leveled (i.e., solutions to problems at a level depend only on solutions to problems at the previous level), the formulation is called **serial**, else it is called **non-serial**.

## **DP Categories: {monadic, polyadic} × {serial, non-serial}**

---

- ♦ Based on these two criteria, we can classify DP formulations into four categories
  - serial-monadic (e.g., 0/1 knapsack)
  - non-serial-monadic (e.g., longest common subsequence)
  - serial-polyadic (e.g., Floyd's all-pair shortest paths)
  - non-serial-polyadic (e.g., optimal matrix parenthesization)
- ♦ This classification is useful since it identifies concurrency and dependencies that guide parallel formulations.

# a Concrete Example: Smith-Waterman

## ◆ Linear gap penalty

$$H_{i,j} = \max \left\{ \begin{array}{l} 0 \\ H_{i-1,j} + g \\ H_{i,j-1} + g \\ H_{i-1,j-1} + W(a_i, b_j) \end{array} \right\}$$

## ◆ Affine gap penalty

$$E_{i,j} = \max \left\{ \begin{array}{l} E_{i,j-1} - G_{\text{ext}} \\ H_{i,j-1} - G_{\text{init}} \end{array} \right\}$$

$$F_{i,j} = \max \left\{ \begin{array}{l} F_{i-1,j} - G_{\text{ext}} \\ H_{i-1,j} - G_{\text{init}} \end{array} \right\}$$

$$H_{i,j} = \max \left\{ \begin{array}{l} 0 \\ E_{i,j} \\ F_{i,j} \\ H_{i-1,j-1} + W(a_i, b_j) \end{array} \right\}$$

# **Algorithmic Tricks/Improvements**

- ◆ Search space reduction
- ◆ Query profile (uses texture cache)
- ◆ Data layouts
- ◆ Input-size dependent choice of algorithms
- ◆ Speculation
- ◆ Storage reduction
- ◆ Processing order
- ◆ Fine-tuning block/thread counts
- ◆ SIMD instructions
- ◆ Use of CPU and GPU
- ◆ Use of local memory
- ◆ Multi-GPU
- ◆ Calculation time prediction equation
- ◆ Pipelining
- ◆ Recompile GPU code on the fly
- ◆ Use of BWT
- ◆ Tiling
- ◆ Block pruning

# **LTDP & Rank Convergence**

- ♦ LTDP definition
  - its solution and the solutions of its subproblems are in the domain of the tropical semiring
  - the subproblems can be grouped into a sequence of stages, such that the solution to a subproblem in a stage depends on only the solutions in the previous stage
  - this dependence is linear in the tropical semiring
- ♦ Examples
  - 1) Longest Common Subsequence; 2) Viterbi
  - 3) Needleman-Wunsch; 4) Smith-Waterman
- ♦ Empirical rank convergence help parallelization

\* S. Maleki, M. Musuvathi, and T. Mytkowicz, “Parallelizing dynamic programming through rank convergence,” ACM SIGPLAN Notice, 2014.

\* S. Maleki, M. Musuvathi, and T. Mytkowicz, “Efficient parallelization using rank convergence in dynamic programming algorithms,” CACM, 2016.

# **Topic Overview**

---

- ♦ Overview of Serial Dynamic Programming
- ♦ Serial Monadic DP Formulations
- ♦ Nonserial Monadic DP Formulations
- ♦ Serial Polyadic DP Formulations
- ♦ Nonserial Polyadic DP Formulations

# **Serial Monadic DP Formulations**

---

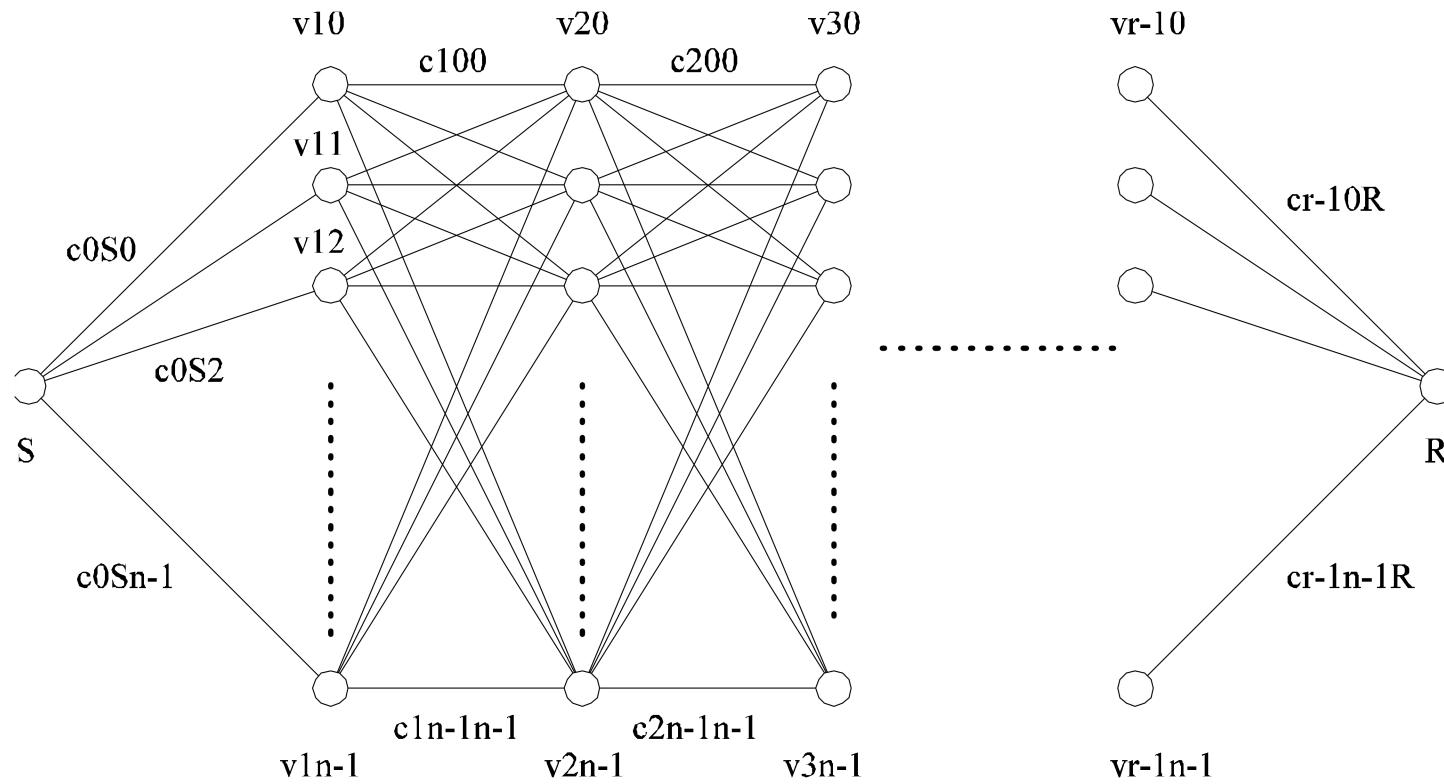
- ♦ It is difficult to derive canonical parallel formulations for the entire class of formulations.
- ♦ For this reason, we select two representative examples, the shortest-path problem for a multistage graph and the 0/1 knapsack problem.
- ♦ We derive parallel formulations for these problems and identify common principles guiding design within the class.

# **Shortest Path for a Multistage Graph**

- ♦ Special class of shortest path problem where the graph is a weighted multistage graph of  $r + 1$  levels.
- ♦ Each level is assumed to have  $n$  levels and every node at level  $i$  is connected to every node at level  $i + 1$ .
- ♦ Levels zero and  $r$  contain only one node, the source and destination nodes, respectively.
- ♦ The objective of this problem is to find the shortest path from  $S$  to  $R$ .

# **Shortest Path for a Multistage Graph**

- ◆ An example of a serial monadic DP formulation for finding the shortest path in a graph whose nodes can be organized into levels.



# **Shortest Path for a Multistage Graph**

- ♦ The  $i$ -th node at level  $l$  in the graph is labeled  $v_i^l$  and the cost of an edge connecting  $v_i^l$  to node  $v_j^{l+1}$  is labeled  $c_{i,j}^{l+1}$ .
- ♦ The cost of reaching the goal node  $R$  from any node  $v_i^l$  is represented by  $C_i^l$ .
- ♦ If there are  $n$  nodes at level  $l$ , the vector  $[C_0^l, C_1^l, \dots, C_{n-1}^l]^T$  is referred to as  $C^l$ . Note that  $C^0 = [C_0^0]$ .
- ♦ We have  $C_i^l = \min \{(c_{i,j}^{l+1} + C_j^{l+1}) \mid j \text{ is a node at level } l+1\}$

# **Shortest Path for a Multistage Graph**

- ♦ Since all nodes  $v_j^{r-1}$  have only one edge connecting them to the goal node  $R$  at level  $r$ ,
- ♦ the cost  $C_j^{r-1}$  is equal to  $c_{j,R}^{r-1}$ .
- ♦ We have:

$$C^{r-1} = [c_{0,R}^{r-1}, c_{1,R}^{r-1}, \dots, c_{n-1,R}^{r-1}].$$

Notice that this problem is serial and monadic.

# **Shortest Path for a Multistage Graph**

- ♦ The cost of reaching the goal node  $R$  from any node at level  $l$  is ( $0 < l < r - 1$ ) is

$$C_0^l = \min\{(c_{0,0}^l + C_0^{l+1}), (c_{0,1}^l + C_1^{l+1}), \dots, (c_{0,n-1}^l + C_{n-1}^{l+1})\},$$

$$C_1^l = \min\{(c_{1,0}^l + C_0^{l+1}), (c_{1,1}^l + C_1^{l+1}), \dots, (c_{1,n-1}^l + C_{n-1}^{l+1})\},$$

⋮

$$C_{n-1}^l = \min\{(c_{n-1,0}^l + C_0^{l+1}), (c_{n-1,1}^l + C_1^{l+1}), \dots, (c_{n-1,n-1}^l + C_{n-1}^{l+1})\}.$$

# **Shortest Path for a Multistage Graph**

- ♦ We can express the solution to the problem as a modified sequence of matrix-vector products.
- ♦ Replacing the addition operation by minimization and the multiplication operation by addition, the preceding set of equations becomes:

$$\mathcal{C}^l = M_{l,l+1} \times \mathcal{C}^{l+1},$$

where  $\mathbf{C}^l$  and  $\mathbf{C}^{l+1}$  are  $n \times 1$  vectors representing the cost of reaching the goal node from each node at levels  $l$  and  $l + 1$ .

# **Shortest Path for a Multistage Graph**

- ♦ Matrix  $M_{l,l+1}$  is an  $n \times n$  matrix in which entry  $(i, j)$  stores the cost of the edge connecting node  $i$  at level  $l$  to node  $j$  at level  $l + 1$ .

$$M_{l,l+1} = \begin{bmatrix} c_{0,0}^l & c_{0,1}^l & \cdots & c_{0,n-1}^l \\ c_{1,0}^l & c_{1,1}^l & \cdots & c_{1,n-1}^l \\ \vdots & \vdots & & \vdots \\ c_{n-1,0}^l & c_{n-1,1}^l & \cdots & c_{n-1,n-1}^l \end{bmatrix}.$$

- ♦ The shortest path problem has been formulated as a sequence of  $r$  matrix-vector products.

# **Parallel SP for a Multistage Graph**

- ♦ We can parallelize this algorithm using the parallel algorithms for the matrix-vector product.
- ♦  $\Theta(n)$  processing elements can compute each vector  $C^l$  in time  $\Theta(n)$  and solve the entire problem in time  $\Theta(rn)$ .
- ♦ In many instances of this problem, the matrix  $M$  may be sparse. For such problems, it is highly desirable to use sparse matrix techniques.

# 0/1 Knapsack Problem

- ◆ We are given a knapsack of capacity  $c$  and a set of  $n$  objects numbered  $1, 2, \dots, n$ . Each object  $i$  has weight  $w_i$  and profit  $p_i$ .
- ◆ Let  $v = [v_1, v_2, \dots, v_n]$  be a solution vector in which  $v_i = 0$  if object  $i$  is not in the knapsack, and  $v_i = 1$  if it is in the knapsack.
- ◆ The goal is to find a subset of objects to put into the knapsack so that

$$\sum_{i=1}^n w_i v_i \leq c$$

(that is, the objects fit into the knapsack) and

$$\sum_{i=1}^n p_i v_i$$

is maximized (that is, the profit is maximized).

# 0/1 Knapsack Problem

- Let  $F[i, x]$  be the maximum profit for a knapsack of capacity  $x$  using only objects  $\{1, 2, \dots, i\}$ . The DP formulation is:

$$F[i, x] = \begin{cases} 0 & x \geq 0, i = 0 \\ -\infty & x < 0, i = 0 \\ \max\{F[i - 1, x], (F[i - 1, x - w_i] + p_i)\} & 1 \leq i \leq n \end{cases}$$

# 0/1 Knapsack Problem

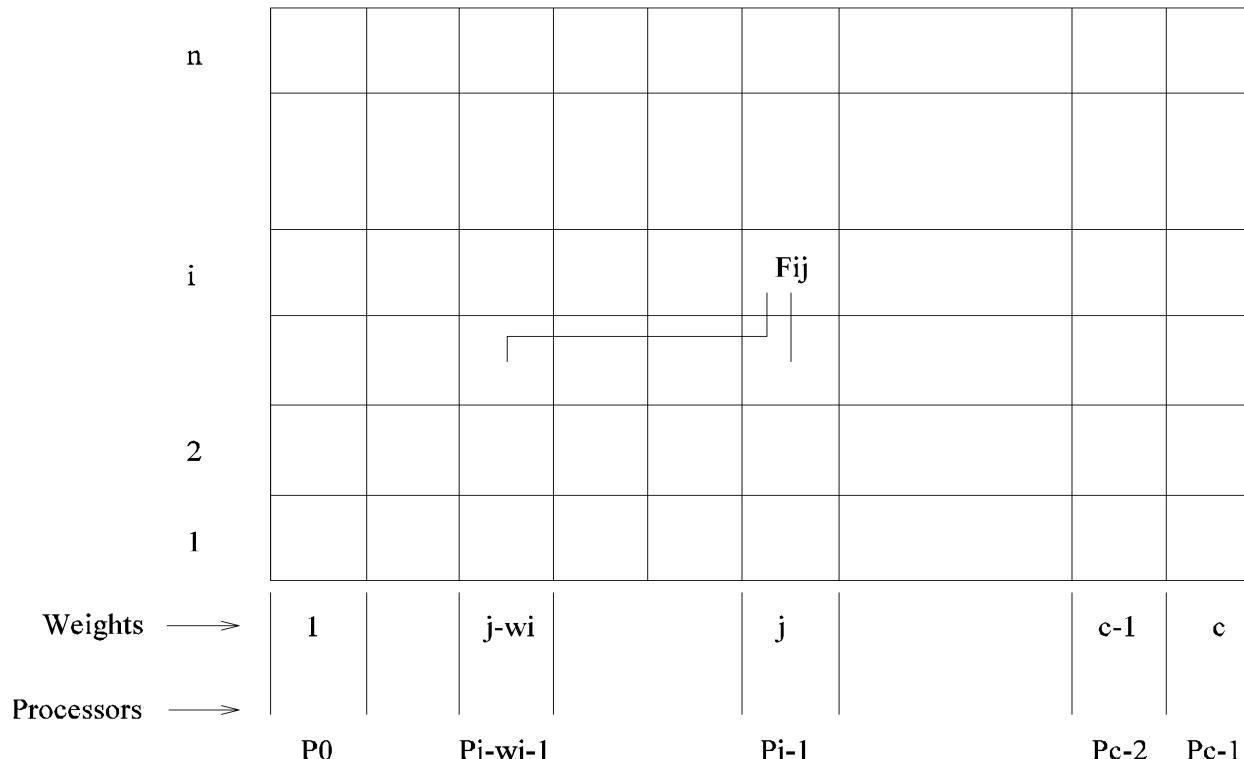
---

- ♦ Construct a table  $F$  of size  $n \times c$  in row-major order.
- ♦ Filling an entry in a row requires two entries from the previous row: one from the same column and one from the column offset by the weight of the object corresponding to the row.
- ♦ Computing each entry takes constant time; the sequential run time of this algorithm is  $\Theta(nc)$ .
- ♦ The formulation is serial-monadic.

# 0/1 Knapsack Problem

- Computing entries of table  $F$  for the 0/1 knapsack problem. The computation of entry  $F[i,j]$  requires communication with processing elements containing entries  $F[i-1,j]$  and  $F[i-1,j-w_i]$ .

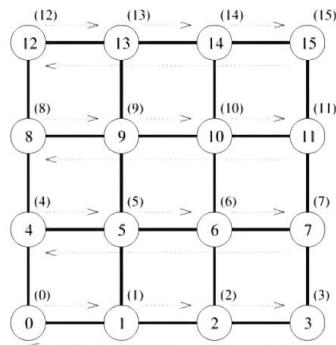
Table F



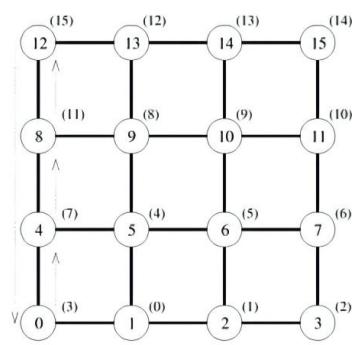
# 0/1 Knapsack Problem: Circular q-Shift

◆ Ring:  $(t_s + mt_w) \min(q, p-q)$

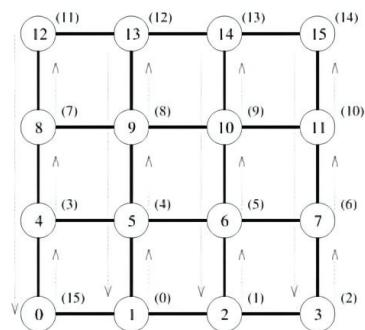
◆ Mesh:  $(t_s + mt_w) (\sqrt{q} + 1)$



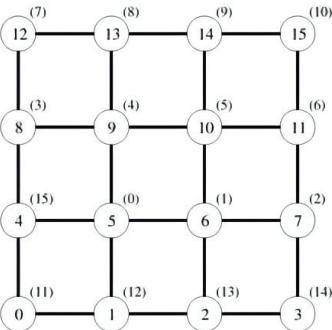
(a) Initial data distribution and the first communication step



(b) Step to compensate for backward row shifts



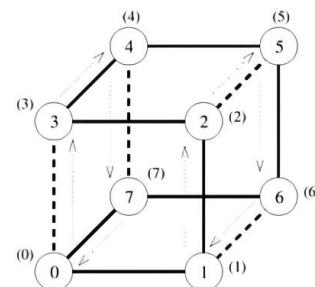
(c) Column shifts in the third communication step



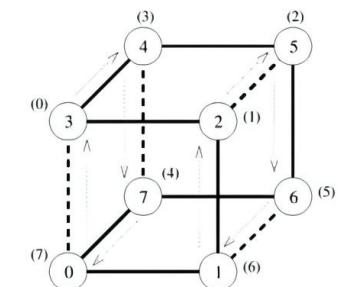
(d) Final distribution of the data

◆ Hypercube:  $(t_s + mt_w) \log p$

- for long msg with bi-directional channels (E-cube routing):  $t_s + m t_w$

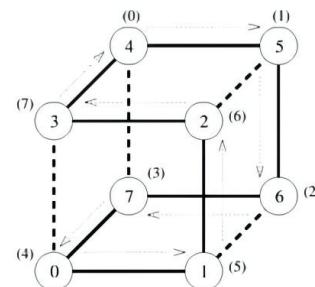


First communication step of the 4-shift

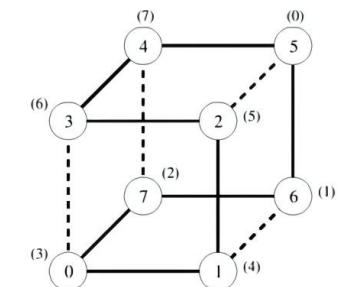


Second communication step of the 4-shift

(a) The first phase (a 4-shift)



(b) The second phase (a 1-shift)



(c) Final data distribution after the 5-shift

examples of the 5-shift on a mesh and a hypercube

# 0/1 Knapsack Problem

- ♦ Using  $c$  processors in a PRAM, we can derive a simple parallel algorithm that runs in  $\underline{O(n)}$  time by partitioning the columns across processors.
- ♦ In a distributed memory machine, in the  $j$ -th iteration, for computing  $F[j,r]$  at processing element  $P_{r-1}$ ,  $F[j-1,r]$  is available locally but  $F[j-1,r-w_j]$  must be fetched.
- ♦ The communication operation is a circular shift and the time is given by  $(t_s + t_w) \log c$ . The total time is therefore  $t_c + (t_s + t_w) \log c$ .
- ♦ Across all  $n$  iterations (rows), the parallel time is  $\underline{O(n \log c)}$ . Note that this is not cost optimal.

# 0/1 Knapsack Problem

- ♦ Using  $p$ -processing elements, each processing element computes  $c/p$  elements of the table in each iteration.
- ♦ The corresponding shift operation takes time  $(2t_s + t_w c/p)$ , since the data block may be partitioned across two processors, but the total volume of data is  $c/p$ .
- ♦ The corresponding parallel time is  $n(t_c c/p + 2t_s + t_w c/p)$ , or  $O(nc/p)$  (which is cost-optimal).
- ♦ Note that there is an upper bound on the efficiency of this formulation:  $1 / (1 + t_w/t_c)$

## **Nonserial Monadic DP Formulations: Longest-Common-Subsequence**

- ♦ Given a sequence  $A = \langle a_1, a_2, \dots, a_n \rangle$ , a subsequence of  $A$  can be formed by deleting some entries from  $A$ .
- ♦ Given two sequences  $A = \langle a_1, a_2, \dots, a_n \rangle$  and  $B = \langle b_1, b_2, \dots, b_m \rangle$ , find the longest sequence that is a subsequence of both  $A$  and  $B$ .
- ♦ If  $A = \langle c, a, d, b, r, z \rangle$  and  $B = \langle a, s, b, z \rangle$ , the longest common subsequence of  $A$  and  $B$  is  $\langle a, b, z \rangle$ .

# Longest-Common-Subsequence Problem

- ♦ Let  $F[i,j]$  denote the length of the longest common subsequence of the first  $i$  elements of  $A$  and the first  $j$  elements of  $B$ . The objective of the LCS problem is to find  $F[n,m]$ .
- ♦ We can write:

$$F[i, j] = \begin{cases} 0 & \text{if } i = 0 \text{ or } j = 0 \\ F[i - 1, j - 1] + 1 & \text{if } i, j > 0 \text{ and } x_i = y_j \\ \max \{F[i, j - 1], F[i - 1, j]\} & \text{if } i, j > 0 \text{ and } x_i \neq y_j \end{cases}$$

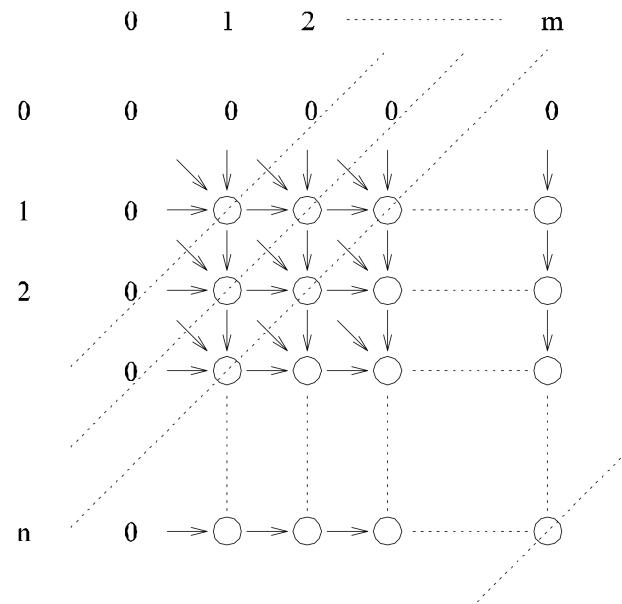
# **Longest-Common-Subsequence Problem**

---

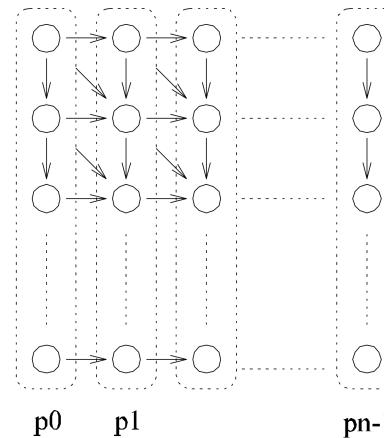
- ♦ The algorithm computes the two-dimensional  $F$  table in a row- or column-major fashion. The complexity is  $\Theta(nm)$ .
- ♦ Treating nodes along a diagonal as belonging to one level, each node depends on two subproblems at the preceding level and one subproblem two levels prior.
- ♦ This DP formulation is nonserial monadic.

# Longest-Common-Subsequence Problem

- (a) Computing entries of table for the longest-common-subsequence problem. Computation proceeds along the dotted diagonal lines.
- (b) Mapping elements of the table to processing elements.



(a)



(b)

# Longest-Common-Subsequence: Example

- Consider the LCS of two amino-acid sequences **H E A G A W G H E E** and **P A W H E A E**. For the interested reader, the names of the corresponding amino-acids are A: Alanine, E: Glutamic acid, G: Glycine, H: Histidine, P: Proline, and W: Tryptophan.

	H	E	A	G	A	W	G	H	E	E
P	0	0	0	0	0	0	0	0	0	0
A	0	0	0	1	1	1	1	1	1	1
W	0	0	0	1	1	1	2	2	2	2
H	0	1	1	1	1	1	2	2	3	3
E	0	1	2	2	2	2	2	3	4	4
A	0	1	2	3	3	3	3	3	4	4
E	0	1	2	3	3	3	3	3	4	5

- The F table for computing the LCS of the sequences. The LCS is **A W H E E**.

# Parallel Longest-Common-Subsequence

---

- ◆ Table entries are computed in a diagonal sweep from the top-left to the bottom-right corner.
- ◆ Using  $n$  processors in a PRAM, each entry in a diagonal can be computed in constant time.
- ◆ For two sequences of length  $n$ , there are  $2n-1$  diagonals.
- ◆ The parallel run time is  $\Theta(n)$  and the algorithm is cost-optimal.

# Parallel Longest-Common-Subsequence

- ◆ Consider a (logical) linear array of processors. Processing element  $P_i$  is responsible for the  $(i+1)$ -th column of the table.
- ◆ To compute  $F[i,j]$ , processing element  $P_{j-1}$  may need either  $F[i-1,j-1]$  or  $F[i,j-1]$  from the processing element to its left. This communication takes time  $t_s + t_w$ .
- ◆ The computation takes constant time ( $t_c$ ).
- ◆ We have:  $T_P = (2n - 1)(t_s + t_w + t_c)$ .
- ◆ Note that this formulation is cost-optimal, however, its *efficiency* is upper-bounded by 0.5!
- ◆ Can you think of how to fix this?

# **Serial Polyadic DP Formulation: Floyd's All-Pairs Shortest Path**

- Given weighted graph  $G(V, E)$ , Floyd's algorithm determines the cost  $d_{i,j}$  of the shortest path between each pair of nodes in  $V$ .
- Let  $d_{i,j}^k$  be the minimum cost of a path from node  $i$  to node  $j$ , using only nodes  $v_0, v_1, \dots, v_{k-1}$ .
- We have:

$$d_{i,j}^k = \begin{cases} c_{i,j} & k = 0 \\ \min \{d_{i,j}^{k-1}, (d_{i,k}^{k-1} + d_{k,j}^{k-1})\} & 0 \leq k \leq n - 1 \end{cases}.$$

- Each iteration requires time  $\Theta(n^2)$  and the overall run time of the sequential algorithm is  $\Theta(n^3)$ .

# **Serial Polyadic DP Formulation: Floyd's All-Pairs Shortest Path**

- ♦ A PRAM formulation of this algorithm uses  $n^2$  processors in a logical 2D mesh. Processor  $P_{i,j}$  computes the value of  $d_{i,j}^k$  for  $k=1,2,\dots,n$  in constant time.
- ♦ The parallel runtime is  $\Theta(n)$  and it is cost-optimal.
- ♦ The algorithm can easily be adapted to practical architectures

# **Nonserial Polyadic DP Formulation: Optimal Matrix-Parenthesization Problem**

- ♦ When multiplying a sequence of matrices, the order of multiplication significantly impacts operation count.
- ♦ Let  $C[i,j]$  be the optimal cost of multiplying the matrices  $A_i, \dots, A_j$ .
- ♦ The chain of matrices can be expressed as a product of two smaller chains,  $A_i, A_{i+1}, \dots, A_k$  and  $A_{k+1}, \dots, A_j$ .
- ♦ The chain  $A_i, A_{i+1}, \dots, A_k$  results in a matrix of dimensions  $r_{i-1} \times r_k$ , and the chain  $A_{k+1}, \dots, A_j$  results in a matrix of dimensions  $r_k \times r_j$ .
- ♦ The cost of multiplying these two matrices is  $r_{i-1}r_kr_j$ .

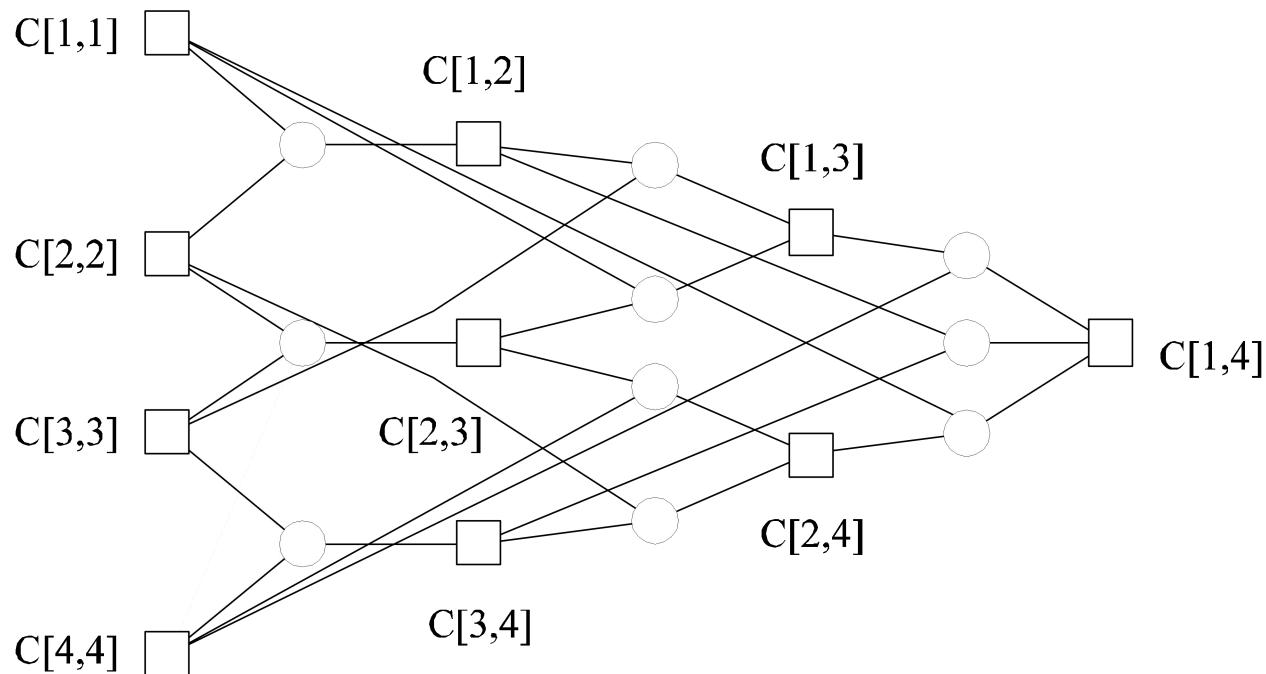
# **Optimal Matrix-Parenthesization Problem**

- ♦ We have:

$$C[i, j] = \begin{cases} \min_{i \leq k < j} \{C[i, k] + C[k + 1, j] + r_{i-1}r_kr_j\} & 1 \leq i < j \leq n \\ 0 & j = i, 0 < i \leq n \end{cases}$$

# **Optimal Matrix-Parenthesization Problem**

- ♦ A nonserial polyadic DP formulation for finding an optimal matrix parenthesization for a chain of four matrices. A square node represents the optimal cost of multiplying a matrix chain. A circle node represents a possible parenthesization.



# **Optimal Matrix-Parenthesization Problem**

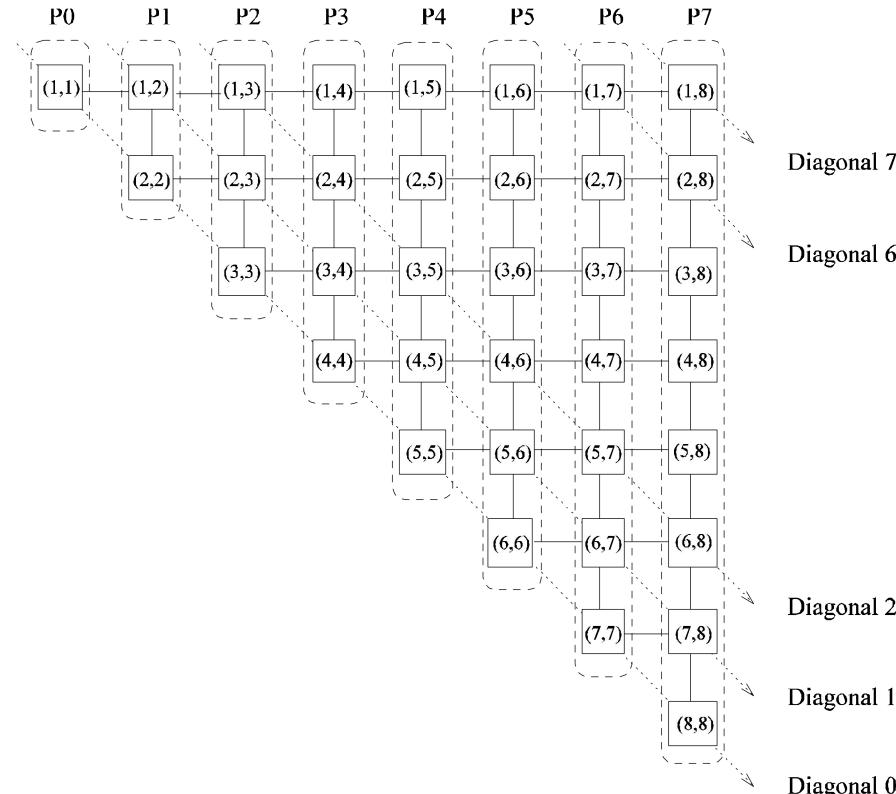
- ♦ The goal of finding  $C[1,n]$  is accomplished in a bottom-up fashion.
- ♦ Visualize this by thinking of filling in the  $C$  table diagonally. Entries in diagonal  $l$  corresponds to the cost of multiplying matrix chains of length  $l+1$ .
- ♦ The value of  $C[i,j]$  is computed as  $\min\{C[i,k] + C[k+1,j] + r_{i-1}r_kr_j\}$ , where  $k$  can take values from  $i$  to  $j-1$ .
- ♦ Computing  $C[i,j]$  requires that we evaluate  $(j-i)$  terms and select their minimum.
- ♦ The computation of each term takes time  $t_c$ , and the computation of  $C[i,j]$  takes time  $(j-i)t_c$ . Each entry in diagonal  $l$  can be computed in time  $lt_c$ .

# **Optimal Matrix-Parenthesization Problem**

- ♦ The algorithm computes  $(n-1)$  chains of length two. This takes time  $(n-1)t_c$ ; computing  $n-2$  chains of length three takes time  $(n-2)t_c$ . In the final step, the algorithm computes one chain of length  $n$  in time  $(n-1)t_c$ .
- ♦ It follows that the serial time is  $\Theta(n^3)$ .

# *Optimal Matrix-Parenthesization Problem*

- ♦ The diagonal order of computation for the optimal matrix-parenthesization problem.



# Parallel Optimal Matrix-Parenthesization Problem

- ◆ Consider a logical ring of processors. In step 1, each processor computes a single element belonging to the 1-th diagonal.
- ◆ On computing the assigned value of the element in table C, each processor sends its value to all other processors using an all-to-all broadcast.
- ◆ The next value can then be computed locally.
- ◆ The total time required to compute the entries along diagonal 1 is  $lt_c + t_s \log n + t_w(n-1)$ .
- ◆ The corresponding parallel time is given by:

$$\begin{aligned} T_P &= \sum_{l=1}^{n-1} (lt_c + t_s \log n + t_w(n-1)), \\ &= \frac{(n-1)(n)}{2}t_c + t_s(n-1) \log n + t_w(n-1)^2. \end{aligned}$$

# Parallel Optimal Matrix-Parenthesization Problem

- When using  $p$  ( $< n$ ) processors, each processor stores  $n/p$  nodes.
- The time taken for all-to-all broadcast of  $n/p$  words is

$$t_s \log p + t_w n(p - 1)/p \approx t_s \log p + t_w n$$

and the time to compute  $n/p$  entries of the table in the  $l$ -th diagonal is  $l t_c n/p$ .

$$\begin{aligned} T_P &= \sum_{l=1}^{n-1} (l t_c n/p + t_s \log p + t_w n), \\ &= \frac{n^2(n-1)}{2p} t_c + t_s(n-1) \log p + t_w n(n-1). \end{aligned}$$

$$T_P = \Theta(n^3/p) + \Theta(n^2)$$

- This formulation can be improved to use up to  $n(n+1)/2$  processors using pipelining.
  - due to the non-serial nature of the problem

# *Discussion of Parallel Dynamic Programming Algorithms*

- ♦ By representing computation as a graph, we identify three sources of parallelism: i) parallelism within nodes, ii) parallelism across nodes at a level, and iii) pipelining nodes across multiple levels. The first two are available in serial formulations and the third one in non-serial formulations.
- ♦ Data locality is critical for performance. Different DP formulations, by the very nature of the problem instance, have different degrees of locality.

# Summary

---

- ♦ Several types of Dynamic Programming
  - Serial Monadic DP Formulations
    - example: shortest path for a multistage graph
    - example: 0/1 knapsack problem
  - Nonserial Monadic DP Formulations
    - example: longest-common-subsequence problem
  - Serial Polyadic DP Formulations
    - example: all-pairs shortest path
  - Nonserial Polyadic DP Formulations
    - example: optimal matrix-parenthesization
- ♦ Sources of parallelism
  - (serial) within a subproblem; and among intra-level subproblems
  - (nonserial) pipelining